An Auditable Agentic Architecture for Regime-Aware Credit Trading Simulation Across Rating-Segmented Universes*

Alejandro Reynoso

Abstract

This paper presents a reorganized and substantially expanded architecture for governance-first credit trading simulation across AAA, AA, and A rating universes. The revised manuscript is structured around four major sections, each divided into four fully developed subsections, so that the conceptual argument, system design, operational logic, and institutional implications are explained with greater depth and continuity. The architecture remains explicitly non-live and research oriented. Its purpose is not to promise deployable alpha, but to demonstrate how a regime-aware credit research process can be decomposed into intelligible modules that preserve rating discipline, surface evidence, document intermediate reasoning, constrain capital allocation deterministically, and produce auditable decision records.

At the center of the design is a hybrid division of labor. Deterministic components generate synthetic spread paths, engineer structured evidence, infer market regimes, calculate exposures, apply execution and liquidity haircuts, run adverse scenarios, and preserve the final audit bundle. LLM-enabled agents operate only within those boundaries. Rating-bucket agents analyze bounded credit evidence for the AAA, AA, and A universes; a coordinator synthesizes cross-bucket posture; and a writer agent converts the artifact trail into a committee-ready narrative. In this arrangement, the language model is not a discretionary trader. It is a governed analytical component embedded inside a system whose authority ultimately rests on deterministic controls, explicit schemas, and human review.

The expanded paper argues that the institutional value of agentic finance lies less in theatrical autonomy than in disciplined decomposition. Credit strategies differ by instrument, liquidity, basis exposure, rating migration sensitivity, and stress behavior. A credible architecture must therefore explain what is being simulated, why rating segmentation matters, how regime inference shapes posture, what controls override model enthusiasm, which artifacts are preserved, and where human responsibility remains non-delegable. The resulting framework is offered as a research demonstrator, a teaching instrument, and a committee-reviewable prototype for auditable credit allocation workflows.

1 Strategic Framing of Regime-Aware Credit Trading

1.1 Why Credit Trading Requires Taxonomy Before Architecture

Credit trading is often introduced through a deceptively simple image: an investor buys a bond when spreads appear attractive, holds the instrument while income accrues, and later sells when the market reassesses risk more optimistically. That description is not wrong, but it is radically incomplete. In institutional practice, credit trading is a family of related but materially different activities that only partly overlap in economic logic, required data, time horizon, balance-sheet

*A companion Google Colab notebook for this paper is available at https://colab.research.google.com/drive/17R_96UPkNplHkjnrdKIIbXC0ZD3htqEd?usp=sharing.

usage, and governance burden. Cash corporate bond carry trades, long-short relative-value positions, credit default swap hedges, rating-migration trades, curve positions, index overlays, capital-structure arbitrage, and liquidity-defense trades all live inside the broad credit universe, but they do not rely on the same evidence or fail for the same reasons. When a paper proposes an agentic architecture for credit trading, the first responsibility is therefore not software design but conceptual discipline. The architecture has to state what kind of credit trading is being represented, what risks are being abstracted away, and why the chosen simplifications remain economically meaningful rather than merely convenient. Classical credit-risk research has already shown that the pricing of debt, the shape of spreads, and the interpretation of default risk depend on structure, term, recovery assumptions, and market conditions rather than on a single generic notion of "credit" alone. A serious architecture begins where that lesson is acknowledged rather than ignored. [2, 3, 4]

That conceptual discipline matters even more when the system is described as agentic. If the designer starts from software roles instead of market taxonomy, there is a strong temptation to invent agents that sound sophisticated but do not correspond to real institutional functions. One agent may be told to "find opportunities," another to "manage risk," and a third to "write recommendations," yet such labels remain shallow if the underlying market object has not been defined. Credit markets do not present a unified decision problem. A trader in a cash bond book asks about spread compensation, funding, liquidity, and benchmark fit. A trader in synthetic credit asks in addition about contract standardization, counterparty exposure, basis risk, collateral mechanics, and settlement conventions. A portfolio manager rotating between rating buckets thinks about relative spread compensation, downgrade sensitivity, capacity, turnover, and committee tolerance for temporary mark-to-market pain. Each of these functions calls for distinct evidence and distinct control logic. The reason taxonomy must come before architecture is therefore straightforward: without it, the designer cannot justify why one agent exists, what it is allowed to see, how its output should be constrained, or which deterministic checks must sit above it.

The present paper addresses that problem by focusing on one bounded domain inside the larger credit landscape: rating-segmented spread trading across synthetic AAA, AA, and A universes. This is not a claim that investment-grade credit is the whole market. It is an explicit narrowing of scope so that the proposed architecture can preserve economic coherence. In these universes, spread behavior is shaped not only by the probability of outright default but also by liquidity conditions, rate volatility, balance-sheet capacity, benchmark demand, sector rotation, and the market's changing willingness to be paid modest premia for assuming modest but nontrivial credit uncertainty. That makes investment-grade credit an especially useful setting for a governance-first research demonstrator. The instruments are similar enough that rating segmentation is interpretable, yet different enough that a regime-aware model can show why capital posture should not be static. In calm markets, the system may rationally tolerate more exposure to the A bucket. In widening regimes, it may prefer capital preservation, migration toward higher quality, and tighter execution assumptions. Those decisions are meaningful precisely because the architecture has been built around a clear object: not all credit trading, but a specific, bounded family of spread-sensitive allocation decisions. [5, 6]

The emphasis on taxonomy also clarifies the relationship between human institutions and machine components. Committee processes inside real asset managers are not organized around vague abstractions such as "the AI thinks spreads look good." They are organized around roles, evidence packs, risk limits, and documented disagreements. A credit committee wants to know which part of the book is expected to contribute return, which part is held primarily for quality defense, what regime assumptions underlie the recommendation, how much turnover the trade implies, and what scenario would most quickly invalidate the thesis. The architecture proposed here reflects that reality by treating software roles as proxies for institutional functions: rating-bucket analysts

interpret evidence inside bounded universes, a coordinator consolidates posture across buckets, deterministic modules enforce limits and compute exposures, and a reporting layer translates the run into committee language. By deriving agents from institutional roles rather than from generic "AI" categories, the model becomes easier to audit, easier to critique, and easier to teach.

Another reason taxonomy matters is that the relevant measure of success changes with the trading object. A directional high-yield thesis might reasonably be judged by realized carry versus default loss. A market-neutral relative-value trade might be judged by spread convergence net of funding and liquidity costs. A hedge may be considered successful even if it loses money in benign markets, provided it preserves capital in adverse states. The architecture in this paper is designed around research questions that fit rating-segmented investment-grade spread trading: can the system distinguish quality buckets, infer regime-sensitive posture, explain why one bucket deserves more or less capital, apply execution and liquidity realism, and preserve a complete record of how each recommendation emerged? Returns still matter, but they are not sufficient. What matters equally is whether the chain of reasoning is institutionally legible. That is why the architecture is described as auditable before it is described as intelligent. It aims to show how agentic tools may support disciplined credit research without collapsing the distinction between explanation, recommendation, and authority. [10, 11, 12]

The section that follows builds on this framing by explaining why rating segmentation is not merely a stylistic choice but the organizing principle of the simulation. Once the paper has established that credit trading is plural rather than singular, the architecture can be evaluated on a more honest basis. The correct question is no longer whether the system can "trade credit" in some broad and theatrical sense. The correct question is whether a bounded, rating-aware, regime-sensitive research process can be represented in a way that is analytically rich, operationally disciplined, and transparent enough for committee review. That narrower question is also the more demanding one, because it forces every component of the architecture to justify its existence in market terms rather than in software rhetoric.

1.2 The Economic Logic of Rating-Segmented Credit Universes

The choice to organize the simulation around AAA, AA, and A universes serves three purposes at once: it aligns the model with recognizable investment-grade practice, it creates a natural bridge between market economics and agent responsibilities, and it establishes a tractable environment for explaining regime-dependent capital posture. Rating buckets are not perfect summaries of credit quality, and no serious credit professional would substitute a letter grade for full fundamental analysis. Yet rating categories remain institutionally powerful because they influence benchmark construction, mandate language, investor expectations, liquidity tiers, and the permissible risk envelope of many portfolios. A governance-first demonstrator should therefore use a segmentation logic that committees already understand. The point is not to say that all AAA credits are interchangeable or that all A credits behave identically. The point is to start from a classification that has real organizational meaning, then show how a regime-aware process can refine but not discard that structure.

Within the bounded universe used here, each rating cohort plays a distinct analytical role. The AAA bucket functions as the quality anchor. Spread levels are typically lower, outright default expectations are remote, and price behavior is often dominated by macro duration demand, scarcity, benchmark preference, and liquidity conditions rather than by acute idiosyncratic distress. The AA bucket serves as an intermediate layer in which quality remains strong but compensation for risk is less compressed and sensitivity to macro repricing can become more visible. The A bucket, although still investment grade, usually provides the richest spread premium among the three and

therefore absorbs a larger share of cyclical optimism and cyclical anxiety. When the market is stable, A-rated exposure may offer a more attractive balance between incremental carry and acceptable risk. When the market moves toward stress, that same bucket often expresses widening first and most visibly. Structuring the simulation around these three universes allows the model to ask an intelligible question: how should capital migrate when the trade-off between spread income and quality defense changes through time? [1, 2, 6]

This segmentation also creates an elegant connection between economics and architecture. Because each rating bucket has a different expected behavior, each bucket agent can be given a bounded interpretive task rather than an unbounded trading mandate. The AAA agent can be asked to focus on quality preservation, liquidity resilience, and whether tight spreads leave too little compensation for duration and execution risk. The AA agent can examine whether the middle of the quality spectrum offers a balanced posture between defense and return. The A agent can assess whether incremental spread compensation remains adequate under the prevailing regime. None of these agents decides the final portfolio on its own. Instead, each produces a structured judgment tied to the economic function of its bucket. This is important for governance. It means the prompt design, evidence bundle, and response schema can all be aligned with a specific segment of the market rather than with a vague request to "analyze credit." The architecture thereby gains precision and the audit trail gains interpretability.

The use of rating buckets also improves the pedagogical value of synthetic data generation. In a purely issuer-level simulation, the designer must invent leverage paths, refinancing schedules, recovery outcomes, and idiosyncratic event processes for many entities at once. That can be informative, but it also risks overwhelming the paper's core contribution, which is governance-first decomposition. By contrast, a rating-bucket design allows synthetic spreads, volatility, liquidity discounts, and turnover costs to be generated at the cohort level while still preserving meaningful dispersion. The paper can therefore demonstrate how regime inference, capital limits, and auditability operate without pretending to have solved issuer-level corporate finance. This is a strength, not a weakness. Research demonstrators are most persuasive when their abstractions are explicit. The present model does not claim to infer the true balance-sheet health of a real issuer. It claims to show how a controlled process can analyze relative posture across recognizable quality bands.

Another advantage of rating segmentation is that it supports a more disciplined treatment of performance evaluation. In many backtests, the temptation is to compress all outcomes into a single return series and judge the model on that basis alone. But rating-aware portfolios are usually evaluated in more nuanced terms. Committees care about where return came from, how much spread risk was assumed to earn it, whether quality drift occurred gradually or abruptly, and how liquidity assumptions changed under stress. A rating-segmented architecture can attribute outcomes more cleanly. If the portfolio outperforms because it leaned into A-rated carry during calm conditions, that fact can be documented. If it preserves capital because the coordinator rotated toward AA and AAA during widening regimes, that can be documented as well. The architecture becomes more than a signal generator; it becomes a mechanism for explaining the relationship between capital posture and market state. Such explanation is especially important in investment-grade credit, where a strategy may appear conservative in name while still accumulating meaningful spread beta beneath the surface. [5, 6]

The broader institutional implication is that rating segmentation provides a common language between researchers, risk managers, and decision committees. A portfolio manager can challenge the simulated allocation by asking why the A bucket deserved incremental weight under a given regime. A risk manager can test whether the assumed liquidity haircut was too mild for AA credits in a stress state. A committee chair can assess whether the proposed posture is consistent with mandate language that emphasizes quality preservation. Those are healthy forms of challenge, and

Table 1: Governance-Oriented Taxonomy for the Bounded Credit Simulation

Trading Modality	Core Object	Primary Decision Question	Treatment in This Paper
Cash bond carry	Investment-grade bond exposure	Is spread compensation sufficient after duration and liquidity adjustment?	Represented in simplified form through synthetic cohort spreads and execution haircuts.
Relative-value quality rotation	Rating-bucket overweights and underweights	Which quality tier offers the best regime-adjusted balance of carry and resilience?	Central object of the architecture.
Synthetic protection overlay	CDS or index hedge	Should defensive credit protection offset widening risk?	Discussed as an extension, not implemented in the baseline model.
Rating migration defense	Exposure shift toward higher quality	When should capital move upward in quality even if carry declines?	Encoded through coordinator posture and deterministic constraints.
Liquidity defense	Trading-cost and capacity discipline	How much apparent alpha survives realistic execution conditions?	Embedded directly in slippage, haircut, and turnover assumptions.
Committee review workflow	Governance and approval process	Can the recommendation be justified, challenged, and archived?	Treated as a first-class output of the system rather than an afterthought.

the architecture is designed to support them. Rather than obscuring its decisions inside a single composite signal, it decomposes exposure into interpretable quality bands. This decomposition does not eliminate uncertainty, but it makes disagreement productive. That is one of the central design goals of the paper.

1.3 Research Scope, Boundaries, and the Value of Explicit Non-Live Positioning

One of the most important design choices in this paper is also one of the least glamorous: the architecture is declared non-live, synthetic, and research oriented from the outset. This is not legal hedging language inserted for convenience after the model has been described. It is a substantive component of the architecture itself. A governance-first system must define the domain within which its outputs are meaningful, and that domain in the present case is limited to committee-reviewable simulation. The model does not connect to brokers, does not source live market data, does not verify the legal terms of securities, does not estimate financing lines, does not manage collateral, and does not autonomously place orders. Those omissions matter because they clarify the type of claim the paper is making. The claim is not that a language-model-enhanced workflow can safely replace an investment team. The claim is that such a workflow can be organized in a way that makes its reasoning inspectable, its limits explicit, and its artifacts durable enough for serious review.

There are practical advantages to that explicit boundary. In live trading, operational and legal complexity quickly overwhelms conceptual clarity. A real implementation must consider data entitlements, stale prints, benchmark methodology, settlement conventions, broker routing, custody constraints, compliance approvals, valuation hierarchy, and the institution's tolerance for model override. Those are essential concerns, but they are not the main subject of this paper. By remaining in a synthetic environment, the design can isolate the governance logic that would

otherwise be buried under operational detail. The architecture can therefore show, with unusual transparency, how evidence is assembled, how regimes are inferred, how LLM outputs are bounded, how deterministic rules constrain exposure, and how the resulting narrative is packaged for human review. Explicit scope limitation thus becomes a methodological advantage. It keeps the paper honest about what has been demonstrated and prevents the reader from confusing a research scaffold with a production-ready trading stack.

This scope discipline is also consistent with modern guidance on model risk, risk data governance, and responsible AI. Supervisory and standards-oriented documents repeatedly emphasize that institutions should define intended use, understand limitations, preserve traceability, and ensure that control functions can challenge the model on understandable grounds. Those expectations become harder, not easier, when a system is allowed to blur the line between analysis and execution. By contrast, a synthetic research architecture can cleanly separate the generation of structured insight from the authority to commit capital. In the present design, LLM components may summarize, compare, and recommend within a fixed schema, but they do not bypass deterministic capital rules and they do not acquire operational agency simply because their prose appears confident. This distinction is crucial. It means the paper treats language models as controlled analytical instruments whose value depends on context, transparency, and bounded responsibility. [10, 11, 12]

The non-live positioning also improves educational usefulness. Many papers in quantitative finance appear sophisticated because they compress a large number of assumptions behind opaque notation or proprietary datasets. That can make the paper look advanced while actually reducing its value as a teaching document. The present manuscript takes the opposite path. It uses synthetic data precisely so that the relationship between assumption, inference, control, and output can be inspected line by line. A reader can see how a regime label changes expected posture, how a liquidity haircut reduces attractive but fragile trades, how the coordinator reconciles disagreements across rating buckets, and how the report cites the artifact trail instead of inventing unsupported claims. This transparency supports teaching, replication, and critique. If a reader disagrees with the regime transition logic or the slippage assumptions, the disagreement can be addressed directly because the architecture makes those assumptions visible.

Another benefit of explicit boundary setting is that it reframes limitations as informative rather than embarrassing. The model does not simulate actual default events, covenant breaches, legal restructuring, recovery waterfalls, or issuer-specific balance-sheet pathologies. It does not claim to estimate a true physical default distribution. It does not incorporate live CDS basis, dealer inventory, or cross-asset contagion from equity volatility and funding markets. These are meaningful omissions, but the right response is not to disguise them. The right response is to state them plainly and explain why the architecture still contributes something worthwhile. What it contributes is a governed pattern for organizing agentic research: data are structured, regimes are inferred, specialized agents interpret bounded evidence, deterministic modules govern exposures, stress scenarios test the posture, and the entire run is converted into an auditable artifact set. That contribution remains valuable even though the model is deliberately not a production engine.

The boundaries of the paper therefore perform a double function. They prevent overclaiming, and they sharpen the institutional lesson. A committee reviewing the model is not asked to trust an opaque machine with capital. It is asked to evaluate whether the decomposition of responsibilities is sensible, whether the audit trail is sufficient, whether the scenario design is serious, and whether the architecture could serve as a disciplined research platform from which more realistic extensions might later be built. In that sense, the explicit non-live posture is not a disclaimer that weakens the paper. It is the condition that makes the paper intellectually responsible.

1.4 Institutional Objectives of the Proposed Demonstrator

Once the market scope and architectural boundaries are clear, the objectives of the demonstrator can be stated precisely. The first objective is interpretive discipline. The system should be able to explain why the spread environment favors or disfavors a given rating bucket without relying on hand-waving language. That requires evidence bundles that are structured enough to support consistent reasoning. Each rating-bucket agent must receive data on current spread levels, recent spread change, volatility, liquidity assumptions, regime state, and simple comparative measures across cohorts. The resulting narrative should not merely say that a bucket "looks attractive." It should say attractive relative to what, under which regime, subject to which execution penalties, and with what downside in the relevant stress case. Interpretive discipline matters because institutional decision makers need language that can be challenged. If a recommendation is not challengeable, it is not governance ready.

The second objective is role separation. LLM components and deterministic components should perform tasks that match their comparative strengths. Deterministic modules are well suited to arithmetic, state tracking, limit enforcement, scenario propagation, turnover measurement, and artifact integrity. They should therefore own those responsibilities completely. Language models are comparatively useful when the task is to compare bounded evidence, synthesize trade-offs, articulate tensions, and translate structured results into committee language. The architecture benefits when these roles are separated rather than blended. A rating-bucket agent can assess whether A-rated spread compensation appears adequate in a neutral regime, but it should not silently modify capital limits or override a liquidity haircut. A writer agent can generate a readable report, but it should not introduce claims unsupported by the stored artifacts. The goal is not to diminish the LLM. The goal is to place it where its contribution is most defensible. [10, 12]

The third objective is governance realism. Many experimental agentic systems treat controls as external commentary that appears after the recommendation has already been formed. In practice, strong governance is formative rather than decorative. Limits, audit requirements, human review, and scenario expectations should shape the decision process from the beginning. That is why the demonstrator embeds controls in the architecture itself. Capital posture must pass deterministic exposure checks. Reported opportunities must survive transaction-cost and liquidity assumptions. Scenario analysis must accompany the recommendation rather than follow it as an optional appendix. Prompt traces, response hashes, model identifiers, and run manifests must exist because committee trust depends on evidentiary continuity. When these elements are integrated, the resulting system no longer resembles a free-form chatbot with a risk memo attached. It resembles a controlled research workflow in which every recommendation is produced under conditions that are legible to oversight functions. [11, 10]

The fourth objective is regime awareness. Credit portfolios do not operate in a stationary environment, and a credit architecture that ignores state dependence will tend to misclassify both opportunity and fragility. Calm markets with stable liquidity and orderly spread compression can reward measured exposure to lower-quality investment-grade cohorts. Stress states can reverse that logic quickly. A research demonstrator should therefore show how regime inference changes posture not only numerically but institutionally. It should show why the same bucket receives different treatment under widening conditions, how execution assumptions become harsher during liquidity stress, and why committee language shifts from opportunity capture to capital preservation when the state changes. Hidden Markov modeling is used here not as a claim of ontological certainty about the market's true regime, but as a disciplined method for imposing state-sensitive structure on a complex time series problem. The value of the regime layer lies in the discipline it imposes on interpretation and control. [7, 8, 9]

The fifth objective is durable documentation. An institutional model is useful only insofar as its outputs survive the moment of generation. If the system recommends a posture but cannot later show the evidence, prompts, control decisions, and scenario results that produced it, then the recommendation has limited governance value. The demonstrator therefore treats artifact preservation as a core deliverable. A run should leave behind an evidence pack, structured agent outputs, coordinator synthesis, stress results, performance tables, and a committee-ready report whose claims can be traced back to stored data. This principle is especially important once LLMs are introduced. Language models can make a workflow more readable, but they can also make it easier to forget how a conclusion emerged. The architecture answers that risk by forcing every narrative layer to remain anchored to deterministic artifacts.

Taken together, these objectives define the meaning of "architecture" in the paper. The architecture is not merely a set of software modules connected in sequence. It is the institutional arrangement through which a bounded credit research question becomes interpretable, state aware, constraint governed, and reviewable. The next section formalizes that arrangement by describing how synthetic evidence is engineered, how regimes are inferred, how rating-bucket agents are bounded, and how the coordinator converts segmented analysis into a unified portfolio posture.

Figure 1: Strategic framing chart.

2 Regime-Aware Multi-Agent Architecture

2.1 Synthetic Spread Data, Feature Engineering, and Evidence Construction

The first operational layer of the architecture is the evidence layer, because no agentic design can be more disciplined than the data it consumes. In a live institutional setting, evidence construction would begin with security masters, benchmark maps, price histories, dealer runs, liquidity scores, spread curves, and mandate metadata. In the present research demonstrator, those objects are represented synthetically so that the mechanics of the architecture can be studied without pretending to solve the much larger problem of production market data engineering. Synthetic generation is not used to escape rigor. It is used to make the assumptions of rigor inspectable. The design therefore begins by specifying a cohort-level data-generating process for the AAA, AA, and A universes in which spread levels, spread changes, volatility, and liquidity penalties evolve through time according to interpretable parameters. Each bucket may have a different baseline spread, a different sensitivity to regime changes, and a different execution penalty. These differences are not decorative. They are the means by which the architecture learns to treat quality tiers differently rather than as relabeled versions of the same object.

The evidence layer is built to reflect the economics of investment-grade spread trading rather than the aesthetics of data abundance. For each cohort and time step, the system constructs a compact but decision-relevant bundle: current spread, recent spread change, rolling volatility, cumulative move over a short window, simple momentum or mean-reversion indicator, liquidity haircut, assumed transaction cost, and a comparison to the other rating buckets. These fields are sufficient to support a serious conversation about posture. If spreads are tight but volatility is rising and liquidity is deteriorating, the evidence should not invite the same conclusion as tight spreads in a quiet and liquid environment. Likewise, an apparently attractive A-rated spread level has a different meaning if AA spreads have widened only modestly than if the entire quality spectrum is repricing violently. The architecture therefore treats context as part of the feature design. A cohort is never analyzed in isolation from the state of its neighbors.

Feature engineering is especially important because agentic systems can easily become verbose without becoming informative. A large prompt filled with loosely related statistics may look sophisticated while actually reducing interpretability. The present design takes the opposite approach. The feature set is intentionally narrow, and each field has a clear institutional purpose. Spread level captures the compensation currently on offer. Spread change captures recent repricing pressure. Volatility measures the stability of that compensation. Liquidity and transaction-cost fields penalize opportunity sets that look attractive only before implementation frictions are recognized. Relative ranking across buckets provides the comparative frame needed for capital rotation decisions. This small set of variables is enough to let the architecture discuss carry, quality, regime sensitivity, and execution realism in a controlled way. It also makes the audit trail cleaner, because the committee can see exactly which facts the agents were allowed to use. [4, 5, 6]

The design of synthetic data also reinforces the paper's distinction between spread trading and default modeling. The generated paths are meant to express plausible cohort-level spread behavior under different market states, not to simulate legal defaults, recovery disputes, or issuer-specific covenant events. That boundary keeps the model honest. A cohort's spread may widen because risk appetite falls, benchmark demand changes, dealer balance sheets tighten, or the regime process shifts toward stress. Those mechanisms matter greatly for investment-grade trading even when actual default remains remote. By centering the evidence bundle on spread and liquidity dynamics, the architecture remains focused on the decision problem it claims to study. It asks whether exposure should rotate across quality buckets when the compensation for spread risk changes, not whether a

specific issuer will file for bankruptcy. The latter question is important, but it belongs to a different modeling layer.

Evidence construction culminates in a structured packet passed to the downstream agents and deterministic components. Each packet contains the current cohort snapshot, a short historical context, cross-bucket comparison statistics, the inferred or estimated regime state, and any precomputed control indicators required by the portfolio layer. Because the packets are machine-readable and stable in schema, they can be hashed, archived, reloaded, and audited. This is a major advantage over ad hoc narrative workflows in which a human analyst copies charts into a slide deck and summarizes them from memory. In the proposed design, the evidence bundle itself is an artifact. If a later review asks why the coordinator reduced A-rated exposure on a given run, the exact evidence packet seen by the A agent can be retrieved and inspected. Such traceability is essential once language-model outputs enter the process.

An equally important property of the evidence layer is that it is extensible without being fragile. The baseline system uses a deliberately concise feature set, but the architecture can absorb more detail if future research demands it. Sector tilts, duration buckets, index-relative measures, financing proxies, CDS overlays, or ESG constraints could be added as new fields without changing the governing principle of the model. What must remain invariant is the discipline with which fields are chosen and documented. Every added variable should have a clear decision role, a clear provenance, and a clear owner inside the control framework. That is the deeper architectural point: evidence engineering is not merely preprocessing. It is the first expression of governance.

2.2 The Hidden Markov Regime Layer and Its Interpretive Role

Once the evidence layer has been defined, the architecture requires a disciplined way to represent market state. Credit spreads are not generated by a stationary mechanism in which the same relationship between spread level, volatility, and liquidity holds in all periods. Markets move between relatively calm states, transitional states, and stress states in which both the meaning of observed spreads and the cost of acting on them change. A regime layer is therefore necessary, not because markets literally announce immutable latent states, but because institutional decision making benefits from an explicit representation of changing conditions. Hidden Markov models provide a useful framework for this task. They allow the system to infer a latent state sequence from observable variables while preserving a transition structure that can be summarized, challenged, and linked to portfolio posture. The use of an HMM in the paper is deliberately pragmatic. It is not presented as the unique or final answer to state detection. It is presented as a disciplined statistical device that turns a noisy spread environment into a manageable state language. [7, 8]

In the baseline design, the regime layer maps the market into three interpretable states: a calm or tightening state, a neutral or mixed state, and a stress or widening state. These labels are chosen because they support institutional discussion rather than because the market naturally labels itself in exactly this way. The calm state is characterized by stable or compressing spreads, moderate volatility, and relatively forgiving liquidity conditions. The neutral state captures transitional environments in which spreads may oscillate without a decisive directional regime. The stress state represents widening pressure, higher volatility, and harsher implementation assumptions. These labels are then attached to posterior probabilities rather than treated as categorical truths. This is an important safeguard. The architecture does not insist that the market is definitively in one state with metaphysical certainty. It records the most probable state and preserves the associated probabilities so that the confidence of the regime inference is visible to the rest of the workflow.

The interpretive value of the regime layer appears in two places. First, it changes how the evidence bundle is read. A 10-basis-point widening in A spreads has a different meaning if it occurs

in a broadly calm environment than if it occurs in a state already characterized by deteriorating liquidity and elevated volatility. Second, it changes how capital governance should respond. Under a calm state, the system may tolerate moderate risk-taking because implementation costs are lower and spread carry is more likely to persist. Under a stress state, the same apparent spread opportunity may be discounted more heavily because turnover becomes costlier and further widening becomes more plausible. The regime layer thus serves as a translation device between observed market data and portfolio posture. It does not choose the trade by itself, but it changes the standards that evidence must meet before risk can be added.

The HMM framework is particularly well suited to a governance-first paper because it produces artifacts that are easy to archive and review. The transition matrix reveals how persistent the states are estimated to be. The filtered and smoothed probabilities show whether the model's conviction is high or low. Regime duration statistics provide an intuitive sense of how often the process moves between calm and stress environments. All of these objects can be saved as part of the run bundle. This makes the state layer more than an internal parameterization; it becomes something the committee can inspect. If a reviewer believes the model is overreacting to noise and entering stress too quickly, the transition statistics offer a concrete basis for challenge. If the model remains in calm too long, that concern can also be expressed quantitatively. The regime layer therefore supports not only better decisions but also better disagreements.

At the same time, the paper is careful not to overstate the power of regime modeling. Hidden Markov methods can help structure the problem, but they do not eliminate model risk. The inferred states depend on the variables included, the sample used, the number of states chosen, and the stability of the underlying process. A governance-ready architecture must therefore treat the HMM as an informative layer subject to review, not as an oracle that immunizes the system from error. This is why regime outputs are paired with deterministic controls and why the final portfolio posture remains subject to committee scrutiny. The statistical model sharpens the conversation; it does not end it.

The larger lesson is that regime awareness becomes institutionally useful only when it is connected to action rules and documentation. In this architecture, the HMM is not included to provide a decorative state label in a chart title. It exists so that evidence packets, agent prompts, execution assumptions, stress scenarios, and coordinator decisions can all be conditioned on a shared representation of market state. That shared state language is one of the main mechanisms through which the architecture achieves coherence across modules. [9, 7, 8]

2.3 Rating-Bucket Agents and the Discipline of Structured LLM Judgment

The agent layer begins only after the evidence and regime layers have already imposed structure on the problem. This sequencing is intentional. An LLM is most defensible in a financial workflow when it receives bounded evidence, performs a bounded interpretive task, and returns a bounded output that can be validated automatically. The architecture therefore assigns one credit agent to each rating bucket: AAA, AA, and A. Each agent receives a standardized packet describing the current state of its cohort, the relative position of that cohort versus the others, the inferred regime, and any control-relevant metrics that the downstream coordinator must consider. The prompt instructs the agent to evaluate attractiveness, fragility, key supporting evidence, main downside risk, and a recommended qualitative posture such as overweight, neutral, or underweight. It is also instructed to justify every conclusion only with the fields present in the packet. The agent is not allowed to invent macro narratives, reference live issuers, or claim knowledge beyond the stored artifacts.

This design choice reflects a broader principle: language models should interpret constrained

evidence, not manufacture the world they are supposed to analyze. In unconstrained settings, LLMs are prone to filling missing context with plausible but unverified narrative. That behavior is precisely what a governance-first research architecture must prevent. The present system does so by making prompt design, schema enforcement, and output validation part of the architecture rather than afterthoughts. Every rating-bucket response is required to fit a machine-readable schema containing fields such as summary thesis, confidence band, evidence bullets, key risk, and recommended posture. Responses that fail schema validation are rejected or repaired before they can affect the rest of the workflow. Because the outputs are structured, they can be compared across runs, hashed for provenance, and analyzed for consistency. This means the LLM layer becomes governable in a way that free-form prose rarely is. [10, 12]

The substantive role of the rating-bucket agents is interpretive specialization. The AAA agent is expected to emphasize capital preservation, liquidity resilience, and the problem of low compensation in very tight spread environments. The AA agent is expected to focus on balance: whether the middle tier offers enough incremental spread without giving up too much quality defense. The A agent is expected to analyze the conditions under which additional spread income remains attractive after accounting for regime and execution fragility. By distributing the task this way, the architecture mirrors a common institutional pattern: specialists analyze segments, while synthesis occurs one layer above. This segmentation has two advantages. First, it reduces cognitive load inside each prompt, which generally improves clarity and consistency. Second, it yields a more interpretable disagreement structure. If the A agent is constructive while the AAA agent remains defensive, the coordinator can explain that the system is not internally inconsistent; it is expressing the trade-off between return seeking and quality preservation.

Operational discipline extends beyond prompt wording. Each agent invocation is logged with a timestamp, model identifier, prompt hash, response hash, token diagnostics if available, and parse status. These metadata are not glamorous, but they are central to the paper's claim of auditability. They allow the reviewer to know exactly which model version was used, whether the response changed across reruns, and whether any post-processing occurred between generation and storage. In addition, the architecture can compare the agent's qualitative posture with deterministic evidence. If the A agent recommends overweight while volatility, liquidity penalties, and regime stress flags all point in the opposite direction, the discrepancy is surfaced rather than silently absorbed. That comparison does not make the agent useless. It makes its contribution contestable, which is precisely what governance requires.

It is also important to note what the rating-bucket agents do not do. They do not compute weights, execute trades, alter transition probabilities, or bypass capital rules. They do not add private information to the workflow. They do not "own" the final portfolio. Their authority is intentionally partial. This partial authority is a design strength because it preserves the distinction between explanation and control. The architecture can therefore benefit from the LLM's comparative advantage in synthesis and natural-language articulation without granting it operational powers that would be difficult to defend institutionally. In a serious research setting, that separation is not a constraint on innovation; it is a condition for innovation that can be reviewed and trusted.

The choice of a compact model class, such as a Claude Haiku-class model, fits this philosophy. The task assigned to the agents is not expansive world modeling but disciplined interpretation of a limited evidence set. A smaller, faster model operating under strict schemas, deterministic checks, and comprehensive logs can be more appropriate than a larger model with more rhetorical range but weaker operational control. The paper's architecture is therefore agnostic about frontier spectacle and focused instead on bounded usefulness.

2.4 Coordinator Workflow, Decision Arbitration, and Committee Narrative Generation

The coordinator is the layer at which segmented judgment becomes portfolio posture. Its job is not simply to average the opinions of the AAA, AA, and A agents. Its job is to reconcile their views in light of regime state, deterministic risk constraints, and the institutional priorities of the portfolio. This distinction matters. If the architecture merely aggregated sentiment scores, it would still be vulnerable to the illusion that language-model confidence is the same as portfolio authority. Instead, the coordinator receives structured outputs from the bucket agents, cross-bucket evidence metrics, current and target exposure states, and any constraint flags from the deterministic control layer. It then produces a synthesis that explains which cohort should be emphasized, which cohort should be de-emphasized, what the overall risk posture should be, and what uncertainties remain unresolved. Because the coordinator operates on structured inputs, its own output can also be schema bound and archived. [11, 10]

The arbitration logic embedded in the coordinator is designed to make disagreement informative. Suppose the A agent argues that spreads now compensate sufficiently for risk, the AA agent recommends balance, and the AAA agent remains defensive because liquidity has weakened. A weak architecture might flatten those views into a single weighted average and lose the economic meaning of the disagreement. The present design instead preserves the tension. The coordinator is expected to state why incremental carry in the A bucket may be attractive, why the regime still justifies caution, and how the final posture balances those competing considerations. This produces a more realistic portfolio conversation. Credit committees rarely face unanimous evidence. They face partial signals that must be reconciled under constraint. The coordinator exists to make that reconciliation explicit rather than implicit.

Yet even the coordinator does not own the final capital decision completely. After producing a recommended posture, the architecture passes the synthesis through deterministic rules that apply maximum exposure limits, turnover caps, liquidity haircuts, and any bucket-level restrictions defined by the research design. If the coordinator wants to overweight the A bucket but the implied turnover breaches a control threshold, the recommendation must be clipped or rejected. If the coordinator proposes a balanced posture but the regime layer is in high-conviction stress, the deterministic overlay may force a more defensive allocation. This is a central feature of the architecture. Narrative intelligence may organize the argument, but deterministic controls retain final authority over positions. That is the only way to ensure that the system's most articulate output is not automatically its most powerful one.

The final narrative layer is handled by a writer agent that receives only the stored artifacts of the run: regime summary, evidence tables, validated bucket outputs, coordinator synthesis, deterministic control decisions, stress-test results, and ex post performance or scenario statistics. The writer's task is to translate this bundle into committee-ready prose. It may summarize the market state, explain the rationale for the proposed posture, identify the principal risks, and describe the scenarios under which the recommendation would be reconsidered. What it may not do is invent supporting evidence or backfill undocumented reasoning. This constraint is critical. A report that is eloquent but not artifact grounded may improve presentation while weakening governance. The architecture therefore treats the writer as a reporting interface to the evidence trail, not as a source of independent analysis.

The result is an end-to-end workflow in which every stage has a distinct institutional meaning. Evidence engineering defines what may be known. Regime modeling defines how market state is represented. Rating-bucket agents interpret bounded evidence. The coordinator synthesizes segmented judgments into a coherent posture. Deterministic rules enforce limits and implementation

Table 2: End-to-End Architecture and Artifact Flow

Layer	Primary Input	Primary Output	Governance Function
Synthetic evidence engine	Cohort assumptions and regime-sensitive parameters	Structured spread and liquidity packets	Makes assumptions explicit and reproducible.
Regime layer	Spread changes, volatility, liquidity signals	State label, probabilities, transition statistics	Conditions interpretation on a documented market state.
AAA/AA/A agents	Cohort evidence packet and regime context	Schema-bound analytical judgment	Produces bounded interpretation tied to a specific rating function.
Coordinator	Validated agent outputs and control flags	Cross-bucket posture recommendation	Preserves disagreement while producing a unified portfolio narrative.
Deterministic allocator	Recommended posture and hard limits	Feasible target weights and turnover check	Retains final authority over exposure and implementation.
Writer and archive layer	Stored artifacts from the full run	Committee report and audit bundle	Converts the run into a durable and reviewable record.

realism. The writer translates the validated result into a document suitable for human scrutiny. This chain is the practical expression of the paper's core claim: agentic value in finance comes not from maximal autonomy but from disciplined decomposition.

One further benefit of this layered workflow is that it gives uncertainty an institutional address. Instead of disappearing into a single portfolio number, uncertainty is distributed across observable locations: low-conviction regime probabilities, disagreement between rating-bucket agents, tight control headroom, and stress scenarios that expose asymmetric downside. The coordinator can therefore communicate not only a target posture but also the quality of the case supporting that posture. This is valuable for committees because a well-governed recommendation should be explicit about whether it is strong but capacity constrained, tentative but directionally useful, or internally divided in a way that warrants caution. By preserving these distinctions, the architecture avoids the false neatness that often makes automated systems look decisive while actually making them harder to trust.

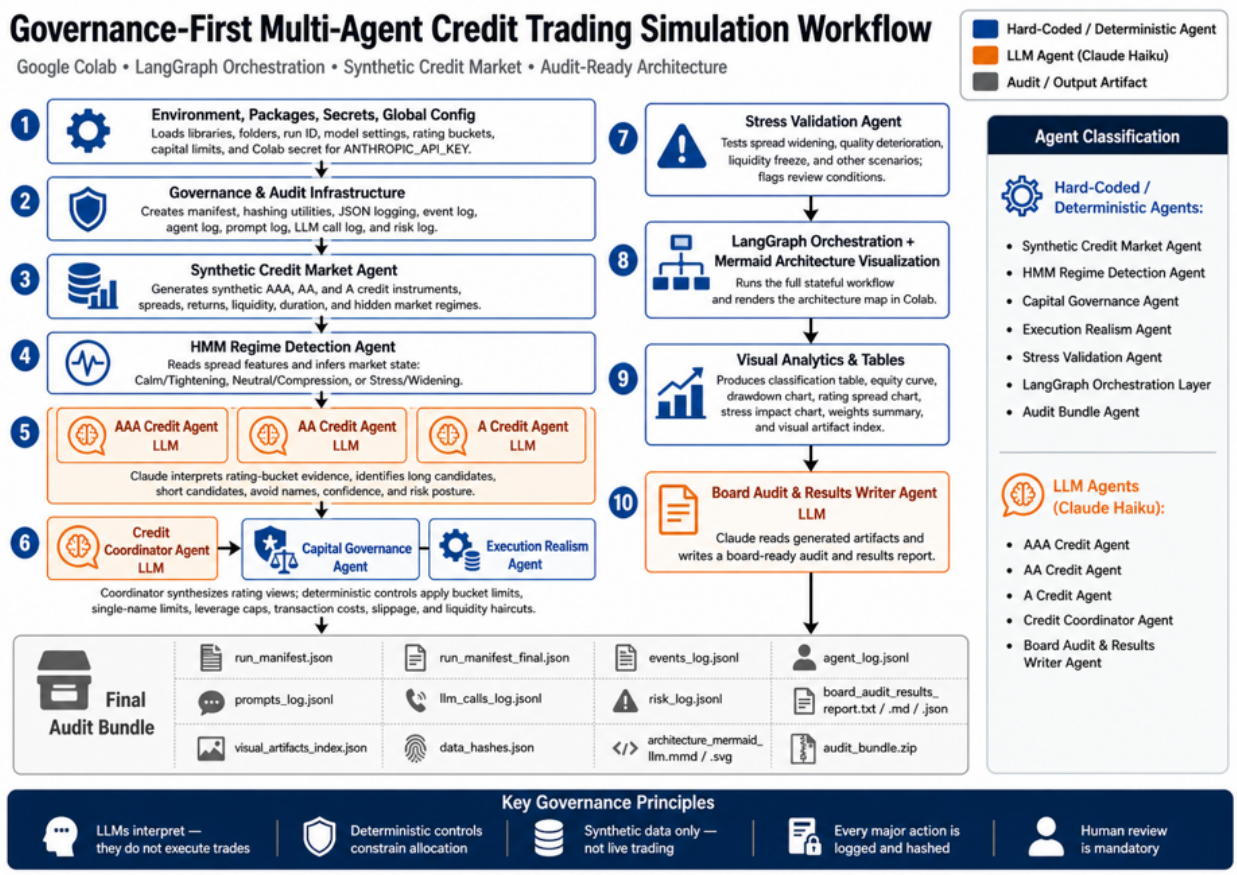

Figure 2: Architecture chart.

3 Governance, Execution, and Audit Design

3.1 Deterministic Controls and the Non-Negotiable Primacy of Capital Governance

The defining feature of a governance-first architecture is not that it contains a risk section. It is that risk discipline determines what the rest of the system is allowed to mean. In the proposed model, deterministic controls are therefore not appended to the recommendation stage as a final cautionary gesture. They are positioned as the binding authority over capital posture. This principle is essential once LLM components are introduced. A language model can summarize evidence persuasively, but persuasive language has no natural stopping point. If the architecture does not contain a separate authority structure capable of saying "no" in arithmetic terms, then the most articulate component will gradually become the most powerful component. The design in this paper prevents that drift by giving deterministic logic final responsibility for exposure limits, turnover bounds, quality floors, liquidity penalties, and implementation feasibility. In practical terms, the agent layer may propose; the deterministic layer disposes.

This priority is implemented through a hierarchy of controls. At the portfolio level, the architecture maintains gross and net exposure bounds, minimum and maximum allocations by rating bucket, optional regime-specific caps, and turnover limits that prevent the model from appearing skillful only by trading unrealistically often. At the cohort level, the control stack can impose tighter constraints in stressed regimes, such as reduced maximum weight for the A bucket or a required minimum allocation to higher quality cohorts when state probabilities indicate elevated widening risk. At the workflow level, the system validates schema compliance, rejects malformed agent outputs, checks for missing evidence fields, and falls back to conservative defaults if required inputs are unavailable. These checks are simple by design. Their purpose is not to mimic every detail of an enterprise risk platform. Their purpose is to ensure that every path from evidence to allocation passes through explicit, reviewable rules.

Capital governance also has a conceptual role: it clarifies what sort of authority the architecture is permitted to exercise. The model is allowed to organize information, compare trade-offs, and recommend posture inside a bounded research environment. It is not allowed to redefine the institution's risk appetite through prose. That boundary reflects a broader lesson from model-risk guidance: a decision system becomes safer not when it produces more confident outputs, but when the relationship between recommendation and approval remains visible and contestable. The present design expresses that lesson operationally. Exposure proposals are clipped if they exceed limits. Regime stress can trigger a defensive overlay even if the narrative layer remains constructive. Liquidity penalties can reduce the attractiveness of trades that look profitable in frictionless terms. Human review remains the final authority above the deterministic layer. This is the architecture's answer to the problem of automation theater. It does not confuse a well-written recommendation with governed permission. [10, 11]

Another reason deterministic controls must lead is that they stabilize the meaning of backtest and simulation results. A strategy that appears strong without turnover limits, liquidity haircuts, or quality floors may not be strong once those conditions are imposed. By embedding controls in the run itself rather than in an ex post caveat, the architecture ensures that reported outcomes already reflect the discipline under which the strategy claims to operate. This is crucial for committee credibility. It is one thing to present a favorable simulated return path and later mention that implementation would have been harder in reality. It is another thing, and a much more institutional thing, to show performance only after implementation constraints have been enforced. The latter is what the demonstrator aims to do.

The control layer also supports constructive challenge. Because each rule is explicit, a reviewer can debate it directly. If a committee believes the A-bucket cap in stress is too restrictive, that can be discussed as a parameter choice. If the turnover ceiling is too lenient for the liquidity assumptions, that can be revised transparently. Such discussions are healthy because they focus on observable governance design rather than on the opaque internals of a language model. In other words, deterministic controls move disagreement to the level where institutions are strongest: policy, limits, assumptions, and accountable choices.

The broader implication is that capital governance is not merely compatible with agentic research; it is the condition that makes agentic research institutionally meaningful. Without hard limits, structured overrides, and conservative fallbacks, the architecture would be an interesting narrative engine. With them, it becomes a bounded decision process whose claims can be evaluated in terms relevant to real oversight.

This is also why escalation rules matter. If multiple control thresholds are breached, if schema repair fails repeatedly, or if the regime layer signals stress while the narrative layer remains aggressively constructive, the correct architectural response is not to improvise a more persuasive justification. It is to escalate, flag the run, and require human review under a clearly conservative default. The presence of such escalation logic is one of the clearest signs that the architecture takes governance seriously.

3.2 Execution Realism, Portfolio Accounting, and Capacity Discipline

Many simulated trading systems appear more intelligent than they are because they operate in frictionless space. They identify attractive relative value, rebalance instantly, and report neat returns without asking whether anyone could have implemented the path at the stated size or under the stated market conditions. Credit trading is especially vulnerable to this distortion because transaction cost, dealer balance-sheet availability, and market depth vary materially across quality buckets and regimes. An architecture that ignores execution realism may therefore misclassify fragile ideas as robust ones. The present paper treats implementation as part of the architecture itself. Every portfolio change is evaluated through estimated transaction costs, slippage, liquidity haircuts, and turnover accounting that worsen when regimes deteriorate. This is not intended as a full market-impact model. It is intended as a disciplined refusal to confuse quoted spread opportunity with realized portfolio opportunity.

Execution realism begins with simple but meaningful assumptions. Each rating bucket is assigned a baseline transaction-cost profile reflecting the fact that higher quality paper may trade with tighter implied execution penalties than lower quality paper under normal conditions, while all cohorts become more expensive to reposition when liquidity worsens. Slippage is then allowed to scale with regime stress so that a spread widening episode not only changes the attractiveness of exposure but also the cost of moving into or out of that exposure. This matters conceptually. In calm periods, rotating toward A-rated carry may be feasible at moderate cost. In stress periods, even a theoretically attractive spread level may be less compelling because the cost of adjustment rises just as conviction becomes most uncertain. The architecture encodes that reality by allowing regime and liquidity assumptions to reduce the net attractiveness of turnover-intensive moves. [5, 6]

Portfolio accounting is handled with equal transparency. The system tracks current weights, proposed target weights, realized turnover, implementation cost, spread-driven mark-to-market effects, and simple carry proxies where appropriate. This accounting layer is deterministic because the institution must be able to explain performance without appealing to narrative interpretation. If the portfolio underperforms, the reviewer should be able to see whether the cause was adverse spread movement, excessive turnover, an overly aggressive tilt toward the A bucket, or repeated

allocation changes during unstable regimes. Such decomposition is essential for learning. A model that only reports total return leaves too many questions unanswered. By contrast, an accounting layer that separates allocation decision from execution cost and market move can show which part of the architecture is adding value and which part is creating fragility.

Capacity discipline follows naturally from this accounting logic. Even in a synthetic environment, the architecture should acknowledge that some decisions become less credible as required turnover or assumed trade size increases. Capacity is not modeled here through a full order book or dealer inventory simulator, but it is represented through the combination of liquidity haircuts, turnover caps, and scenario-specific execution penalties. These mechanisms ensure that the model cannot repeatedly exploit small spread differences as if unlimited balance-sheet capacity were always available. The result is a simulation that favors durable posture changes over hyperactive micro-adjustment. That preference is both realistic and pedagogically useful. It aligns the architecture with the cadence of institutional credit decision making, where large books are often moved deliberately rather than continuously.

Execution realism also improves the quality of human review. A committee can challenge a recommendation more effectively when the implied implementation burden is visible. If the coordinator suggests a sharp rotation from AAA to A during a mixed regime, the committee can ask whether the expected incremental spread truly compensates for the required turnover and reduced liquidity resilience. If the answer depends on unrealistically cheap execution, the trade idea is weaker than it first appeared. By exposing this logic in the accounting layer, the architecture turns implementation from a hidden assumption into an explicit object of governance.

The outcome is a more honest form of simulated intelligence. The system does not merely identify where spreads look rich or cheap. It evaluates whether the portfolio can move into that posture under assumptions that preserve institutional credibility. In a paper concerned with auditability and governance, that distinction is not secondary. It is one of the main reasons the simulation remains useful rather than theatrical.

Equally important, the accounting layer creates a bridge between research and oversight language. Portfolio managers may speak in terms of spread pickup and relative value, while risk officers often focus on turnover, drawdown, and implementation burden. A deterministic accounting record lets both audiences inspect the same recommendation through different but compatible lenses. That shared vocabulary strengthens the committee process because it reduces the chance that a trade appears attractive only by being described in the most flattering metric.

3.3 Stress Testing, Adverse Scenarios, and the Role of Human Committee Review

No credit architecture is credible if it only explains why a trade should work. A governed architecture must also explain how the trade could fail, how quickly losses could accumulate, and what evidence would justify reducing or reversing the posture. Stress testing is therefore treated in this paper not as decoration but as an adversarial companion to recommendation. The system's baseline scenarios are chosen to reflect the core vulnerabilities of rating-segmented investment-grade spread trading: broad parallel spread widening, disproportionate widening in the lower-quality investment-grade bucket, liquidity freeze conditions that amplify turnover cost, and regime transition errors in which the market deteriorates faster than the model recognizes. Each scenario is applied deterministically so that its impact on bucket returns, portfolio drawdown, and posture robustness can be measured and archived. The goal is not to predict the future path of distress with precision. The goal is to expose the recommendation to structured adversity before it is presented to a committee.

This stress discipline has two institutional benefits. First, it improves portfolio logic. A

recommendation that looks attractive in expected terms may still be unacceptable if its loss profile under plausible widening or liquidity shock is inconsistent with mandate tolerance. Second, it improves communication. Committees do not simply ask what the model likes; they ask what could go wrong if the model is wrong. By generating scenario tables and narrative summaries in advance, the architecture makes that question central rather than reactive. If the portfolio is overweight A in a calm regime, the stress module should show what happens under a sudden widening event and whether the quality defense embedded elsewhere in the book is sufficient. If the system is defensive and underweights carry, the stress module should also reveal the opportunity cost of staying conservative through benign spread compression. A good governance process evaluates both risk of action and risk of inaction.

Human committee review is the institutional location where these scenario results gain meaning. The paper is explicit that the model does not self-authorize. The committee remains the forum in which the recommendation, controls, and adverse scenarios are assessed together. This is not a symbolic gesture. Human review introduces forms of challenge that are difficult to encode completely in statistical or prompt logic: concerns about mandate drift, doubts about market microstructure assumptions, skepticism about the persistence of the inferred regime, or knowledge that the institution has broader portfolio exposures not represented in the simulation. A governance-first architecture should welcome such challenge because it is precisely the mechanism through which bounded model outputs are converted into responsible institutional decisions. [10, 11]

The stress module is also useful because it reveals the interaction between model components. A widening scenario does not merely reduce expected return; it may increase turnover costs, invalidate an agent's constructive thesis, and trigger deterministic exposure caps. This interaction matters because real portfolio stress rarely arrives through a single channel. Spread movement, liquidity deterioration, and governance tightening often occur together. By preserving stress outputs as part of the artifact trail, the architecture lets reviewers see how the full chain reacts rather than only how one numerical layer responds. This helps distinguish robust architectures from brittle ones. A brittle system may produce attractive recommendations until the first serious scenario is applied. A robust system shows, in advance, how its own controls and narrative posture change when conditions worsen.

Stress testing further supports learning across runs. Because scenarios are standardized, different policy settings can be compared on a like-for-like basis. A stricter A-bucket cap, a harsher liquidity haircut, or a more defensive regime mapping can be evaluated not only by baseline performance but by how the architecture behaves under repeated adversity. This makes the demonstrator useful as a research platform rather than merely as a single narrative case study. It allows governance design itself to become an object of experimentation.

The key point is simple: in credit research, seriousness is revealed not by how convincingly the architecture describes opportunity, but by how deliberately it preserves the right to challenge itself. Stress testing and committee review are the institutional forms of that self-challenge. [9, 11]

The same logic suggests that scenario design should evolve with the institution's concerns rather than remain static. Over time, committees may decide that downgrade clustering, sector-specific spread shocks, or policy-driven liquidity withdrawals deserve explicit treatment alongside the baseline scenarios. Because the architecture stores scenarios as deterministic artifacts, those additions can be introduced transparently and compared across runs. Stress governance therefore becomes cumulative knowledge rather than ad hoc storytelling.

That cumulative property is one of its main strengths.

Table 3: Control, Stress, and Oversight Matrix

Component	Primary Mechanism	Question It Answers	Why It Matters for Governance
Exposure controls	Bucket caps, quality floors, turnover bounds	Is the recommendation feasible under stated risk appetite?	Prevents narrative enthusiasm from becoming unauthorized capital deployment.
Execution realism	Slippage, liquidity haircut, implementation cost	Does attractive spread compensation survive trading frictions?	Converts paper opportunity into institutionally credible opportunity.
Stress testing	Widening, liquidity shock, regime-error scenarios	How does the posture behave when assumptions deteriorate?	Forces the model to face adverse conditions before committee review.
Schema validation	Parse checks and fallback logic	Can agent outputs be trusted as structured inputs downstream?	Protects the workflow from malformed or uncontrolled LLM output.
Artifact logging	Hashes, manifests, stored packets	Can the full path from evidence to report be reconstructed?	Makes later review and replication possible.
Human committee review	Formal challenge and override authority	Should the institution accept, modify, or reject the posture?	Preserves non-delegable human responsibility.

3.4 Provenance, Logging, Artifact Packaging, and Audit Readiness

Auditability is often described in broad and favorable terms, but its operational meaning is usually vague. In this paper, audit readiness is defined concretely as the ability to reconstruct the full path from assumptions to final report without relying on memory, informal notes, or unverifiable narrative. Achieving that standard requires a provenance layer that is at once mundane and indispensable. Every run should preserve the synthetic data parameters, random seed or generation controls, evidence packets, regime outputs, validated agent responses, prompt and response hashes, coordinator synthesis, deterministic allocation decisions, stress-test tables, and final report text. If charts or summary statistics are generated, they should also be stored in a stable bundle. The result is not merely a folder of outputs; it is an evidentiary chain. That chain is what allows the architecture to claim institutional seriousness rather than merely descriptive sophistication. [11, 10]

The need for provenance becomes more acute, not less, when language models participate in the workflow. Traditional quantitative pipelines can often be rerun from code and data alone, even if the intermediate reasoning remains opaque. LLM-assisted systems add another layer of variability: model versions can change, prompts can be revised, schemas can drift, and textual outputs can differ across otherwise similar runs. Without explicit logging, the institution may be left with an eloquent report but no reliable way to determine exactly how that report was produced. The architecture answers this risk by treating the prompt, model identifier, and parsed response as first-class artifacts. Hashes allow lightweight integrity checking. Stored schemas show what shape the response was required to take. Diagnostics reveal whether the response passed validation or required repair. Together, these records ensure that the LLM layer can be inspected with the same seriousness applied to numerical components. [12, 10]

Artifact packaging also matters because committee communication is not the same thing as system memory. A committee report is designed to be read by humans, not to serve as the sole

technical record of the run. It necessarily compresses information, omits low-level details, and highlights only the most relevant tensions. The architecture therefore maintains a dual output structure. One layer is the committee package: readable prose, summary tables, selected charts, and a concise articulation of posture, risk, and scenarios. The other layer is the technical audit bundle: machine-readable data, logs, manifests, and serialized intermediate objects. This dual structure respects two different institutional needs. Decision makers need clarity. Control functions and researchers need reconstructability. A good architecture serves both without forcing either audience to work from the wrong artifact.

An important secondary benefit of strong provenance is that it improves the quality of model development. When the full chain is preserved, designers can compare runs systematically rather than impressionistically. They can examine whether a change in prompt wording altered the A agent's posture distribution, whether a revised liquidity haircut reduced turnover materially, or whether a different regime mapping produced more stable recommendations. In other words, audit artifacts do not only support oversight after the fact. They support better experimentation during research. This is one reason the paper treats provenance as co-substantial with intelligence rather than as a bureaucratic burden. Documentation is not external to the model's value. It is one of the mechanisms through which that value becomes measurable.

The architecture's audit readiness is completed by the discipline of packaging. At the end of a run, the system assembles a manifest describing the files included, their roles, and their integrity markers. The manifest links the human-readable report to the exact technical artifacts on which it depends. This makes the output portable across stakeholders. A portfolio committee can read the report. A risk function can inspect the stress tables. A model-risk team can review the prompts, hashes, and schemas. A researcher can reproduce the simulation with the stored parameters. The same run therefore supports multiple forms of institutional scrutiny without requiring informal handoff.

For a paper concerned with governed agentic finance, this is the ultimate measure of architecture. A reader should be able to ask not only "what did the system recommend?" but also "what evidence did it use, how was state inferred, what did each specialized agent conclude, which controls overruled or clipped the recommendation, what scenarios were applied, and what permanent artifacts prove the path?" An architecture that can answer all of those questions is much closer to institutional usefulness than one that can merely generate fluent prose about markets.

In that sense, the audit bundle is not the paperwork produced after intelligence has finished its work. It is the durable institutional form that makes the work legible in the first place. A model that cannot preserve this form may still be impressive as a demonstration, but it remains weak as a governed financial instrument.

4 Generative-AI Pathways for Expanding the Agentic Design

This section turns from the baseline governed architecture to a forward-looking research agenda. The objective is not to claim that the present design should immediately be replaced by a more elaborate autonomous stack. The objective is to show how generative AI models could expand the architecture in disciplined stages if the research program chose to pursue greater semantic coverage, richer internal reasoning, adaptive learning, architecture search, and progressively higher autonomy. Each proposed enhancement is therefore framed as an extension path rather than as a declaration that the safer baseline is obsolete. This distinction matters. In institutional finance, the next design is not automatically the better design. A more complex architecture is only superior if it improves analytical breadth, calibration, resilience, and governance together. If it merely adds more agents, more tokens, and more internal conversation without improving the quality of evidence, the clarity of authority, or the durability of the audit trail, then it is not progress but ornament.

The proposals below are written in that governance-first spirit. They ask how generative models might enlarge the analytical field of the system while preserving machine-readable controls, evidence provenance, and committee interpretability. They also ask how the architecture could move from a static, mostly scripted pattern of bounded intelligence toward a more adaptive research platform in which agents learn, criticize one another, search over alternative designs, and eventually approach a higher level of institutional autonomy. For a credit system, these questions are important because the market is information rich, narrative sensitive, and regime dependent. Spreads respond not only to structured balance-sheet variables and synthetic cohort dynamics, but also to management tone, regulation, litigation, policy rumor, supply-chain accidents, macro headlines, dealer commentary, and social-media amplification. A mature generative-AI research agenda should therefore expand coverage beyond the current structured packet without abandoning discipline. The subsections that follow explain one possible path.

4.1 Expanding Agentic Breadth to Ten Specialized Agents for Unstructured Data Intelligence

The current architecture deliberately privileges structured evidence because structured evidence is easier to validate, archive, and govern. That design is appropriate for a first institutional demonstrator. Yet it also leaves analytical territory unexplored. Credit markets do not reprice only because a deterministic spread feature crosses a threshold. They reprice because people revise their beliefs about refinancing conditions, regulation, policy, sector fragility, fraud risk, labor conflict, litigation, geopolitics, management credibility, and investor crowding. Much of that revision process begins in unstructured data long before it is cleanly represented in a benchmark spread series. News headlines, earnings-call transcripts, analyst commentary, blogs, specialized trade press, social-network chatter, and rumor cascades frequently shape the semantic environment in which market participants interpret hard data. If a future version of the architecture aimed to capture a richer representation of that environment, one logical direction would be to expand the agent roster from the current narrow set of rating-bucket specialists to a broader ten-agent design in which several upstream agents are dedicated to unstructured information.

Such an expansion should not be understood as a move toward indiscriminate data hoarding. The purpose of adding agents is not to maximize noise intake. It is to introduce labor specialization into a part of the workflow where labor specialization is genuinely useful. Unstructured data are heterogeneous. A regulatory filing is not the same kind of object as a rumor circulating on a social network. A credible wire-service headline is not the same kind of object as an anonymously sourced market whisper. An earnings-call transcript contains long-form, issuer-linked language, while social

chatter often arrives as fragmented claims, sentiment fragments, memes, and repeated paraphrases. Treating all of these inputs as if they should be processed by a single generic "news agent" would recreate the very conceptual flattening that the baseline paper tried to avoid in the credit domain itself. A more thoughtful extension would allocate different semantic tasks to different agents, then recombine their outputs through a governed evidence-construction process.

In a ten-agent architecture, the first additional agent could be a document-ingestion and triage agent. Its role would be operational rather than judgmental. It would determine whether a piece of unstructured content belongs in the workflow at all, classify it by source type, time sensitivity, and market relevance, and route it to downstream specialist agents. This agent would not decide whether a headline is true or whether a rumor should move capital. It would decide whether the text is about an issuer, a sector, a macro regime, a refinancing channel, or a policy topic relevant to the simulated book. The second additional agent could be a source-reliability and provenance agent. Its task would be to assess how much institutional weight should be assigned to the source before the content is semantically interpreted. Official filings, central-bank statements, recognized financial wires, corporate transcripts, and reputable specialist publications would naturally rank differently from anonymous postings or low-credibility repost chains. This provenance layer is crucial because a generative model is often very good at summarizing a claim even when the claim itself deserves little trust. The architecture therefore needs an explicit mechanism that separates semantic readability from evidentiary credibility.

The third added agent could be a news and event-extraction agent. This agent would convert long-form articles and transcripts into structured event candidates such as guidance revision, legal dispute, management change, spread-relevant policy announcement, funding-market stress indicator, sector disruption, or liquidity event. It would identify named entities, timestamps, directional implications, and any quoted numerical information that could later be checked against structured data. The fourth added agent could be a social-sentiment and crowd-attention agent. Its responsibility would not be to declare that online sentiment is "correct," but to estimate whether attention intensity, directional mood, and narrative concentration around a topic are shifting in ways that may matter for liquidity, volatility, or short-term spread interpretation. The fifth added agent could be a rumor-detection and credibility-staging agent. This one is especially important in credit, because rumor can widen spreads long before facts are resolved. Rather than passing rumor directly into portfolio logic, the architecture would stage rumor into categories such as unverified, partially corroborated, strongly corroborated, or disconfirmed. It would also record the evidence basis for each label so that a later review could distinguish between a market-moving narrative and a high-noise rumor episode.

The sixth agent in the enlarged design could be an entity-resolution and linkage agent. Unstructured texts rarely arrive in a portfolio-ready form. A company may be referenced by legal entity, brand name, ticker shorthand, sector nickname, or the name of a subsidiary. A policy item may affect suppliers, customers, counterparties, or peers indirectly rather than the named firm directly. The role of the linkage agent would be to normalize the referent structure of the incoming text so that event candidates and sentiment scores can be connected to the proper issuer, rating bucket, sector sleeve, and cross-holding map. The seventh agent could be a sector contagion and propagation agent. Its role would be to ask whether a piece of information that appears issuer specific might in fact carry read-through to the wider sector or to adjacent cohorts. A rumored refinery outage, a regulatory enforcement action against a major telecom operator, or a property-market liquidity scare may matter not only for the directly named firm but also for peers and suppliers. In a credit architecture concerned with rating buckets rather than a single issuer, the propagation question is central because semantic shocks frequently travel through sector identity before they become visible in aggregate spread statistics.

The eighth agent could be a macro-narrative alignment agent. One reason unstructured data are difficult to use is that the same headline means different things in different macro states. A hawkish central-bank comment may be interpreted as a growth-positive credibility signal in one environment and as a liquidity-negative tightening signal in another. The macro alignment agent would not replace the statistical regime layer. Instead, it would map semantic developments onto the regime framework already used by the system. It would ask whether a given news flow reinforces the currently inferred regime, introduces contradictory evidence, or suggests a possible regime transition that the statistical layer has not yet absorbed. The ninth agent could be a rating-bucket semantic analyst that specializes in translating the upstream unstructured-data outputs into consequences for AAA, AA, and A cohorts. This agent would act as a bridge between the semantic intelligence layer and the existing rating-bucket agents. It would not produce final weights, but it would frame how news, rumor, sector contagion, and crowd attention alter the interpretive context of each quality band.

The tenth agent should be a governance and evidence-consolidation agent rather than yet another free-form analyst. In a ten-agent system, the greatest danger is not a shortage of insight but a loss of evidentiary hierarchy. Too many semantically eloquent components can easily generate a persuasive fog. The governance-consolidation agent would therefore perform two restrictive functions. First, it would translate the outputs of the preceding unstructured-data agents into a compact, schema-bound packet that can be consumed by the existing coordinator and deterministic control layers. Second, it would quarantine any signal whose provenance, corroboration status, or entity linkage remains uncertain. In other words, the tenth agent would be the guardian of discipline inside a much wider semantic intake. Without that role, an expanded agent roster would risk overwhelming the original architecture with partially checked language.

A ten-agent system of this kind would enlarge agentic breadth in a way that is economically motivated rather than cosmetically impressive. It would let the architecture ask much richer questions than it can ask today. Is widening in the A bucket being driven by structured spread deterioration alone, or is it being accompanied by an emerging negative narrative cluster in sector-specific news? Has a rumor on social networks begun to alter attention even though the underlying claim remains weakly corroborated? Are multiple low-credibility fragments converging on the same refinancing concern, thereby raising the probability that the narrative deserves analyst review even before the market fully reprices? Is a central-bank comment changing the semantic tone of risk in a way that conflicts with the current HMM regime label? These are precisely the kinds of questions that generative models are comparatively good at organizing, provided they are placed in a workflow that disciplines evidence and route selection.

The design challenge lies in how such agents interact. A sensible workflow would be staged. Raw text would first be routed through ingestion and provenance gates. Only then would semantic extraction occur. Rumor and sentiment would be analyzed in parallel rather than sequentially, so that emotional intensity does not automatically become inferred fact. Entity resolution and contagion mapping would then connect those outputs to the credit universe. After that, the macro-alignment agent and rating-bucket semantic analyst would translate text intelligence into portfolio-relevant implications. Finally, the governance-consolidation agent would produce a bounded evidence packet. This staged structure matters because it preserves a crucial institutional principle: text should become portfolio evidence only after it has passed through credibility, linkage, and contextualization filters. A future system that skips these stages might look more "AI native," but it would be less trustworthy.

There are also important governance consequences. Social-network rumors and fast-moving headlines can create procyclical behavior if treated naively. A generative model may summarize a rumor beautifully and still make the architecture more fragile by increasing turnover into noise.

For that reason, any unstructured-data expansion should include explicit controls such as source-weight ceilings, corroboration thresholds, time-decay rules, quarantine states for unresolved claims, and regime-sensitive penalties for rumor-driven repositioning. In a governed credit workflow, the architecture should never be allowed to say, in effect, "social media look worried, therefore increase conviction." It should instead say something more disciplined: "attention around a refinancing rumor has increased, source reliability is low, partial corroboration exists from sector press, entity linkage is high, immediate portfolio consequences remain capped pending confirmation." That kind of phrasing is less dramatic, but it is institutionally superior.

Another issue is model assignment. Not every one of the ten agents needs the same model. Some tasks, such as source classification and entity normalization, may be handled by smaller or fine-tuned specialist models. Others, such as long-context transcript synthesis or contradictory-evidence reconciliation, may justify a more capable generative model. The architecture becomes more efficient when model capacity is matched to task complexity rather than assigned uniformly. This is another reason to think in terms of agent specialization. "More agents" should not mean "more identical LLM calls." It should mean a richer division of analytical labor, with each role engineered for a specific kind of evidentiary transformation.

If implemented carefully, a ten-agent design would not replace the baseline rating-bucket architecture. It would deepen it. The original AAA, AA, and A agents would still evaluate spread, volatility, liquidity, and regime. What would change is the richness of the context they receive. They would no longer consume only synthetic structured features. They would consume a disciplined summary of the semantic field surrounding those features: whether risk narratives are concentrated, whether rumors are intensifying, whether reliable reporting confirms the market's direction, whether sector contagion is plausible, and whether social attention is outrunning evidence. In that sense, unstructured-data expansion is best interpreted not as a separate AI layer glued onto the side of the portfolio process, but as a new evidence frontier that generative models make more tractable.

The broader institutional lesson is that agentic breadth should expand only where additional specialization creates a measurable gain in explanatory power. Credit research genuinely benefits from a deeper reading of unstructured text, but only if the architecture refuses to confuse semantic richness with truth. A ten-agent system is therefore attractive not because it is larger, but because it can divide the difficult labor of reading, filtering, linking, contextualizing, and governing unstructured information. When built in that way, generative AI does not merely make the system more verbose. It makes the evidence space more legible. That is the strongest case for enlarging the roster.

A final benefit of the ten-agent design is methodological clarity. Once the labor of reading unstructured data is decomposed this way, the research team can evaluate each semantic function separately instead of praising or blaming a single monolithic AI layer. One can measure whether provenance scoring reduced false positives, whether rumor staging lowered noise-induced turnover, whether contagion mapping improved sector awareness, and whether the macro-alignment agent added useful regime context. This granularity matters for scientific progress inside the research program. It allows agentic expansion to be justified piece by piece rather than by impression. If the institution discovers that one new agent adds little value, that role can be revised or removed without dismantling the entire design. The result is a more modular path to semantic sophistication. In practical terms, that modularity may be the strongest argument for expanding to ten agents: not that every future deployment must be large, but that a larger, well-partitioned research architecture gives the institution a far better laboratory for understanding which forms of unstructured intelligence genuinely deserve a durable place in credit decision support.

4.2 Reasoning-Rich Multi-Agent Designs: Critique, Adversarial Testing, and Teacher–Student Hierarchies

The baseline architecture is intentionally linear in its analytical logic. Evidence is constructed, a regime is inferred, rating-bucket agents interpret the packet, a coordinator synthesizes posture, deterministic controls clip the recommendation, and a writer communicates the result. That linearity is a virtue at the current stage because it keeps authority easy to locate. Yet linearity also limits what the system can do internally before a recommendation reaches the committee. In particular, it constrains the architecture's ability to challenge its own hypotheses, search for counterarguments, and separate superficially persuasive reasoning from robust reasoning. If a future version of the design wanted to take fuller advantage of generative AI, one of the most important upgrades would be the addition of explicit reasoning architecture: critique agents, adversarial testing agents, counterfactual agents, and teacher–student hierarchies that make internal thought more structured, more testable, and more efficient.

The first conceptual shift is to move from single-pass interpretation to proposal-and-challenge. In many current LLM workflows, an agent is asked to read evidence and produce a recommendation. Even when the prompt includes admonitions to be careful or balanced, the model is still fundamentally engaged in one-pass answer construction. This is often adequate for shallow tasks, but it is weak for institutional analysis because it creates too little organized opposition. Real committees do not operate through one-pass cognition. A credit portfolio idea is normally met by questions about hidden assumptions, implementation fragility, omitted downside cases, source quality, regime instability, and alternative explanations. A more advanced generative architecture should therefore reproduce some of that institutional resistance internally before the recommendation reaches humans. Critique agents are the most direct way to do this.

A critique architecture could begin with a proposer agent that generates the initial portfolio thesis, not unlike the current rating-bucket or coordinator outputs. That thesis would then be passed to one or more critic agents whose job is not to generate a competing recommendation from scratch, but to locate weaknesses in the proposer's argument. One critic might focus on evidence sufficiency. It would ask whether the supporting signals are genuinely strong or whether the conclusion rests on a thin chain of inference. Another critic might focus on data integrity and provenance. It would ask whether the key claims rely on lower-quality sources, unstable entity linkage, or semantically appealing but weakly corroborated narratives. A third critic might focus on portfolio construction logic. It would ask whether the recommendation improves expected posture only by increasing turnover, concentrating exposure, or understating the penalty for stress liquidity. By dividing critique into functions, the architecture prevents "criticism" from degenerating into generic skepticism. Each critic has a domain and a rubric.

The next layer is adversarial testing. A critique agent typically evaluates whether an argument is weak on its own terms. An adversarial agent goes further by deliberately constructing the strongest plausible case against the recommendation. In a credit architecture, this could mean generating the most dangerous alternative explanation for current spread behavior, identifying the macro event that would most quickly invalidate the thesis, or constructing a scenario in which the same evidence should rationally lead to the opposite bucket allocation. The point is not rhetorical combat for its own sake. The point is to expose the recommendation to intelligent opposition while the system still has time to downgrade conviction, tighten limits, or route the case for human escalation. In a governed workflow, adversarial testing is valuable because many fragile trade ideas fail not due to obvious errors, but because nobody forced them to face a better counterargument than they faced in their original prompt.

A mature adversarial layer would likely use several forms of challenge. One form is evidentiary

contradiction. The adversarial agent asks which omitted facts, if true, would reverse the recommendation. Another form is regime contradiction. It asks whether the recommendation is robust if the current regime label is wrong or stale. A third form is narrative contradiction. It asks whether the semantic explanation being offered is merely one plausible story among many and whether a rival story fits the observed data just as well. A fourth form is execution contradiction. It asks whether the trade remains attractive after harsher turnover, liquidity, or timing assumptions are imposed. These different adversarial modes are important because institutional failure rarely comes from only one angle. A trade can be analytically elegant yet operationally impossible; semantically compelling yet weakly evidenced; statistically justified yet vulnerable to a regime shift. A rich reasoning architecture should make these distinct vulnerabilities explicit.

The challenge, of course, is to stop a multi-agent reasoning system from becoming a theatrical debate engine. More agents do not automatically imply better thought. They can also create a long transcript of shallow disagreement. The design must therefore discipline internal reasoning just as strongly as it disciplines external reporting. Each critique or adversarial agent should return structured findings, not open-ended monologues. It should indicate which claim it is challenging, what evidence supports the challenge, how severe the problem is, and whether the issue can be resolved by more data, lower conviction, tighter constraints, or outright rejection. In other words, internal reasoning should produce decision-relevant objects, not merely more text. This is especially important if the architecture is to remain auditable. A committee can review a table of challenged assumptions and resolutions far more effectively than it can review pages of unconstrained AI dialogue.

A second major design pattern is the teacher–student hierarchy. In a financial research workflow, not every case deserves the same amount of reasoning depth or model cost. Many portfolio states are routine. Others are ambiguous, novel, or conflict laden. A teacher–student system uses this asymmetry intelligently. A more capable, more expensive teacher model can be used to design schemas, generate high-quality reasoning exemplars, supervise critique formats, and handle a subset of difficult or low-confidence cases. Smaller student agents can then execute recurring analytical tasks at lower latency and lower cost using those teacher-generated templates or distilled competencies. In this arrangement, the teacher is not simply a bigger chatbot. It is an architectural resource used to upgrade the quality of the entire system.

There are several ways to implement the teacher–student idea. One approach is instructional distillation. The teacher model processes complex historical cases, generates structured rationales and error analyses, and these outputs become a training or templating corpus for smaller student agents. Another approach is escalation-based routing. Student agents handle normal cases until uncertainty, disagreement, or novelty crosses a threshold, at which point the case is escalated to the teacher. A third approach is meta-critique. The teacher does not directly recommend a trade; it reviews whether the students' reasoning architecture has been followed correctly and whether their outputs are internally consistent. Each variant serves a different institutional aim. Distillation improves efficiency, escalation improves case allocation, and meta-critique improves quality control.

Teacher–student hierarchies are particularly attractive in a governed setting because they offer a middle path between brittle hard-coding and expensive full-scale reasoning on every run. The architecture can preserve deterministic controls and modest runtime while still reserving deep semantic reasoning for the moments when it is most needed. In a ten-agent future design, for example, a teacher model might be called only when rumor staging, macro alignment, and rating-bucket interpretation disagree sharply. In a critique-rich architecture, the teacher might adjudicate only when a proposer and multiple critics remain unresolved after a bounded number of rounds. This kind of selective depth is more institutionally plausible than a system that invokes maximum model capacity by default.

A third reasoning enhancement is the introduction of self-explanation without unconstrained private rumination. Financial institutions often want a model to be able to explain itself, but they do not necessarily want to depend on unbounded chains of hidden reasoning that are difficult to validate or govern. A stronger architecture would therefore distinguish between raw model cognition and formalized rationale objects. The system might allow agents to perform internal reasoning, but the only durable artifacts used for governance would be concise argument maps, evidence-to-claim tables, explicit lists of unresolved assumptions, and structured summaries of the most important alternatives considered. This is an important principle for any future reasoning-rich design. The objective is not to maximize visible internal chatter. The objective is to maximize reviewable analytical structure. A critic that says, "three assumptions are unsupported, one downside is untested, and the evidence hierarchy is weak," is far more useful than a critic that produces a page of eloquent but unstructured caution.

The architecture also becomes stronger when critique and adversarial agents are connected to deterministic gates. A critique system should not merely produce objections; it should alter the workflow. If unresolved critique severity is high, the recommendation should be downgraded or forced into human review. If adversarial testing reveals that a small regime misspecification flips the preferred bucket, the system should lower conviction or reduce allowed exposure. If the teacher determines that the students have relied on semantically attractive but weakly evidenced claims, the case should be routed back to evidence assembly. These connections are what convert reasoning from decorative sophistication into portfolio governance. Without them, the architecture risks becoming introspective but not safer.

More advanced versions of the design could include counterfactual agents. These agents would ask what recommendation the system would have made if one key assumption were reversed, if one major source were removed, or if the market were reclassified into an adjacent regime. Counterfactual analysis is valuable in credit because spread interpretation is notoriously sensitive to context. A future counterfactual agent might show, for instance, that the preference for AA over A disappears if liquidity costs normalize, or that the case for defensiveness depends almost entirely on a single macro narrative. Such results make the system's dependencies more legible. They also help the committee distinguish between broad-based conviction and narrow conditional conviction.

Another extension is hierarchical debate. In a hierarchical debate structure, lower-level agents first develop their own findings inside their specialty domains. Mid-level arbiters then resolve contradictions within domain clusters, such as unstructured-data interpretation or portfolio construction critique. Only after that does a senior coordinator assemble the portfolio view. This pattern is superior to a flat debate among many agents because it reduces cross-talk and preserves specialization. It also mirrors how serious institutions reason: domain experts first clarify internal issues, then cross-domain committees address the integrated decision. A future generative architecture that adopted hierarchical debate could be substantially more capable than the baseline while remaining understandable to humans.

The main cost of all these enhancements is complexity. More reasoning layers increase latency, compute expense, observability burden, and the chance of process failure. For that reason, a serious research path should pair every new reasoning component with a termination rule, a cost budget, and a success metric. How many rounds of critique are allowed before the case escalates? Which unresolved conflicts are severe enough to halt the workflow? How much incremental forecast quality or committee confidence does teacher escalation produce relative to cost? How often do adversarial agents correctly identify fragility versus generate false caution? These questions need explicit answers. Otherwise the architecture may become more intellectually entertaining but less operationally coherent.

In the end, the reason to enrich the architecture with critique, adversarial testing, and teacher–

student structures is not that "reasoning" is fashionable. It is that credit decisions are fragile when they are not forced to confront disciplined opposition. A future governed agentic design should therefore aim to institutionalize internal challenge before it increases outward autonomy. That ordering is essential. A system that reasons more deeply but remains governable is a plausible next step. A system that acts more freely without learning how to criticize itself is not. The best advanced architecture is one in which generative models do not merely answer better, but disagree better, challenge better, and teach one another better within a structure the institution can supervise.

There is also an important cultural effect. An institution that adopts reasoning-rich architecture is implicitly teaching itself to value structured dissent rather than theatrical certainty. That cultural shift can be as important as the computational one. When the system routinely presents not just a preferred posture but the strongest objections, the conditions under which those objections matter, and the reason those objections were not fully dispositive, the committee receives a more realistic object for judgment. In turn, human reviewers can respond in kind by challenging the right assumptions rather than reacting to an oversimplified headline conclusion. Over time, the architecture may therefore improve not only model quality but organizational quality. It encourages the research process to become more explicit about uncertainty, opposition, and evidentiary limits. That is precisely the kind of maturity generative AI should support in finance.

4.3 Continuous Learning and the Gradual Transition from Hard-Coded to Adaptive Agents

The baseline architecture described earlier in the manuscript is largely static in its intelligence profile. Its deterministic modules remain fixed unless a human edits them, and its LLM agents behave as bounded interpreters whose knowledge of the current run comes from the evidence packet rather than from accumulated institutional experience. This is an appropriate starting point for a governance-first demonstrator, but it also sets a ceiling on what the system can become. A static agent can be transparent, yet it cannot meaningfully improve except through manual redesign. If the long-term ambition is to create an architecture whose agents become more capable through repeated use, then the next design frontier is continuous learning: a controlled process through which hard-coded or heavily scripted agents gradually acquire better calibration, richer memory, stronger pattern recognition, and more adaptive decision support. The phrase "continuous learning" must be handled carefully in finance, however, because it can mean anything from benign prompt adjustment to unsafe self-modification. The value of this subsection lies in distinguishing those possibilities and outlining a governed path from one to the other.

The first principle of continuous learning is that adaptation should be layered. The system should not jump directly from fixed prompts to self-retraining agents that rewrite their own policies. A much safer progression begins with non-parametric memory. At this first layer, the agent remains the same model with the same core prompt logic, but it gains access to a curated memory store containing prior cases, committee decisions, historical stress episodes, and structured error reports. In practical terms, this means the architecture can retrieve analogues: "what did the committee decide in a similar widening regime when unstructured-data signals were contradictory?" or "which prior cases produced false conviction in the A bucket because liquidity penalties were underestimated?" This memory layer already makes the agent more intelligent in a functional sense, but it does so without changing the model's parameters. It changes context rather than identity. That is often the safest first step toward gradual intelligence.

The second layer is feedback-enriched prompting. Here, the architecture begins to encode lessons from prior runs into explicit policy adjustments. If past evaluations show that the rumor-

staging agent overreacts to low-quality social-network activity, its prompt or scoring rubric can be modified so that corroboration thresholds are stricter. If the coordinator repeatedly underweights execution realism in mixed regimes, the prompt can be updated to require a more explicit turnover justification. This layer remains mostly human authored, but it is increasingly data informed. The system becomes more intelligent because it does not merely run again; it runs again with memory of its prior weaknesses embedded into its instructions. Such a design preserves accountability because every adjustment is explicit, reviewable, and reversible.

The third layer is model-assisted policy tuning. At this stage, generative models themselves can help propose improvements. A meta-learning agent might review a library of past runs, identify recurring error patterns, and suggest revised schemas, new threshold candidates, or new routing rules. Another agent might compare high-performing and low-performing cases to recommend where the architecture is too rigid or too permissive. Importantly, this layer still does not imply autonomous adoption of the proposed changes. The institution can treat these outputs as model-improvement memos subject to human approval. This keeps the benefits of machine pattern detection while preserving oversight. It is also consistent with the broader governance principle that change management in financial models should itself be controlled and documented. [10, 11, 12]

A fourth layer involves learned scoring components that remain subordinate to the architecture's higher-level rules. For instance, the system might develop a learned calibration model that estimates when the sentiment agent is likely to be trustworthy, or a learned routing model that predicts whether a case deserves teacher escalation, or a learned fusion model that combines structured and unstructured signals more effectively than a hand-tuned heuristic. These components make the system more adaptive because they are trained on accumulated history, but they do not yet render the full architecture self-revising. Instead, they behave like learned instruments embedded within a still-governed process. This is an important intermediate stage because it lets the research team gain experience with adaptive behavior in limited domains before permitting broader architectural evolution.

Only after these earlier layers are in place should the design contemplate dynamic training of the agents themselves. Even then, the phrase should refer to scheduled, evaluated, and reversible adaptation rather than constant unsupervised self-updating. A sensible dynamic-training workflow would begin with a curated data-generation pipeline. Every run would produce structured artifacts: evidence packets, predictions, committee decisions, stress-test outcomes, realized errors, overrides, and postmortems. Those artifacts would be turned into training examples only after passing quality filters. False or low-confidence labels, corrupted runs, or cases dominated by unusual operational noise would be excluded or separately tagged. A training steward process would then build supervised datasets for specific learning objectives such as improved event classification, better confidence calibration, more accurate critique severity, or better escalation routing. In this approach, the agent becomes more intelligent over time because it learns from the institution's own history, but the learning remains mediated by a formal pipeline rather than by uncontrolled self-reference.

The next issue is deployment strategy. No continuously learning financial agent should be allowed to replace its predecessor immediately after training. The architecture should instead use champion–challenger testing. The incumbent agent remains the champion. Newly trained variants run in shadow mode on the same evidence, producing recommendations that are logged but not acted upon. Their performance can then be assessed across multiple dimensions: accuracy, calibration, stability under regime shifts, agreement with expert judgment, sensitivity to bad data, and compliance with governance constraints. Only after a challenger consistently outperforms the incumbent on a regime-diverse validation suite should it be promoted. Even then, promotion should be reversible. This design principle is central because learning systems often improve average-case performance while secretly worsening tail behavior. A shadow phase is the simplest way to discover

that before the model acquires operational influence.

If this process is repeated over many cycles, the hard-coded agents begin to change character. They are no longer static role players operating from a one-time prompt design. They become adaptive specialists shaped by the institution's own case history. The sentiment agent becomes better at distinguishing meaningful narrative deterioration from transient noise. The coordinator becomes better at recognizing when disagreements among upstream agents usually signal genuine complexity rather than mere model variance. The critique agents become better at identifying which objections tend to precede real committee overrides. In this sense, the agents do become gradually more intelligent. But the intelligence emerges through controlled accumulation of institutional learning, not through romantic self-awakening. That distinction is important both technically and rhetorically.

A more ambitious extension would be parameter-efficient fine-tuning or adapter-based specialization. Instead of retraining entire models, the architecture could maintain lightweight role-specific adapters updated on curated internal data. One adapter might specialize a general model for rumor credibility staging. Another might specialize the same backbone for critique severity classification. This approach is attractive because it preserves a common model foundation while allowing role differentiation. It also makes rollback easier because adapters can be versioned, tested, and replaced independently. In a large multi-agent architecture, this can be the bridge between generic generative models and deeply institutional agents. It also allows the research program to define "intelligence" more concretely: improved skill on bounded tasks rather than vague claims of general sophistication.

Continuous learning also needs a theory of forgetting. Agents that learn continuously can become brittle if they overweight recent cases and underweight older but structurally important patterns. This is especially dangerous in credit, where the most informative historical episodes may be infrequent stress periods rather than the most recent benign months. A governed learning architecture must therefore maintain replay buffers, regime-balanced evaluation sets, and explicit retention policies so that new learning does not erase old caution. One practical design is to maintain a permanent benchmark library of cases: calm carry environments, rumor-driven widening, policy shocks, liquidity squeezes, sector contagion episodes, and false-positive sentiment waves. Every newly trained agent must pass this library before deployment. In this way, institutional memory is preserved not only in narrative archives but in the validation logic of the evolving system itself.

Another issue is feedback quality. Not every committee decision is a reliable label, and not every realized market outcome proves the model was wrong or right. Markets move for many reasons. Committee choices may reflect mandate considerations outside the model's field of view. If the architecture learns naively from every override or realized spread move, it may internalize noise, politics, or exogenous shocks as if they were clean supervisory signals. Continuous learning therefore needs a feedback taxonomy. Some feedback is hard label material, such as parse failures or objectively wrong entity linkage. Some is soft label material, such as disagreement on conviction or posture. Some is contextual annotation, such as "the committee overrode this recommendation due to external portfolio concentration not represented in the model." A mature learning system will improve faster when it distinguishes these forms rather than flattening them into a single reward signal.

There is also a strategic benefit to continuous learning: it changes the role of human experts. In a static architecture, humans mainly author prompts, choose thresholds, and inspect outputs. In an adaptive architecture, humans also become teachers, curators, and governors of learning. Their expertise is captured in postmortems, correction labels, critique annotations, counterfactual notes, and promotion decisions for challengers. This is a healthier and more realistic human–AI relationship than the fantasy of complete replacement. The institution becomes progressively better at encoding its own judgment into the architecture. That is how hard-coded agents become gradually more

intelligent without ceasing to be institutionally anchored.

The final point is that continuous learning should improve humility before it improves autonomy. A truly mature learning system is not simply one that grows more confident. It is one that becomes better calibrated about when it knows little, when the current case is unlike the training history, when upstream data are contradictory, and when a new regime likely requires human review. If gradual intelligence leads only to more assertive recommendations, the architecture has learned the wrong lesson. The right lesson is deeper: adaptive agents should become better at recognizing complexity, not just better at generating conclusions. In a governance-first credit workflow, that kind of intelligence is far more valuable than superficial autonomy. It is the difference between a model that merely changes over time and a model that matures.

For that reason, the most valuable continuous-learning program may be the one that improves the institution's capacity to remember and compare. Once the architecture stores high-quality case histories, annotated errors, promotion decisions, and regime-balanced evaluation suites, it creates a cumulative research asset that extends beyond any one model generation. New agents can be judged against the memory of older agents. New prompts can be assessed against prior failures. New routing policies can be compared to archived committee outcomes. In effect, the system develops a historical conscience. That conscience is what prevents "learning" from collapsing into drift. It ensures that the architecture becomes gradually more intelligent because it accumulates disciplined experience, not because it is allowed to mutate faster. In a field as path dependent and regime sensitive as credit, this kind of institutional memory is not an optional accessory to adaptive AI. It is the condition under which adaptation remains worthy of trust.

A further implication is that continuous learning should be evaluated not only by immediate improvement in recommendation quality, but also by the stability of institutional language. As agents mature, they should become better at articulating why a case resembles a known pattern, why it departs from precedent, and why those differences matter. This makes learning visible to humans in a way that a hidden gain in accuracy alone does not. Over time, the committee should be able to see that the system is becoming more coherent, better calibrated, and more historically aware. When learning produces that kind of visible improvement, it stops being a technical side project and becomes part of the institution's decision discipline.

4.4 Optimal Model and Workflow Design Through Evolutionary Search and Genetic Algorithms

The baseline architecture, like most serious research systems, is handcrafted. Humans choose how many agents exist, which evidence fields they see, which models they call, in what order they interact, how their outputs are aggregated, and where deterministic rules override narrative judgment. Handcrafting is valuable because it keeps the initial design intelligible. Yet it also means that many architectural choices reflect the bounded imagination of the designers. Why three rating-bucket agents rather than five sector-conditioned agents? Why one coordinator rather than a hierarchy of coordinators? Why a specific set of prompt schemas rather than an alternative decomposition that might produce better calibration? Why a particular routing rule for escalation rather than another? As the agentic design grows more complex, these choices multiply. At some point, manual architecture design becomes its own bottleneck. That is why a future generative-AI research agenda should consider a higher-order optimization layer: the use of genetic algorithms, evolutionary search, or equivalent multi-objective architecture-optimization methods to search for better multi-agent designs.

The case for such methods is strongest when the design space is large, discrete, and only partly understood analytically. Multi-agent architectures have exactly this property. The number of agents,

their specialization boundaries, the topology of their communication, the choice of model family for each role, the prompt template, the use of memory, the escalation thresholds, the routing logic, and the deterministic clipping rules can all be varied. Many of these choices interact nonlinearly. A new critique agent may only add value if a teacher model is also present. A rumor-detection agent may be valuable only when paired with stronger provenance scoring and a harsher quarantine rule. A hierarchical coordinator may outperform a flat coordinator only when the agent roster exceeds a certain size. Because these dependencies are hard to reason about from first principles, it is attractive to treat architecture design itself as a search problem.

A genetic algorithm provides one conceptually clear way to do that. The architecture can be encoded as a genome. One segment of the genome can specify the number and type of agents. Another can specify which evidence sources feed which agents. Another can define the communication topology: linear, star, hierarchical, or debate based. Another can encode model assignment, such as whether a role uses a small specialist model, a general medium model, or a large teacher model. Additional genes can specify critique depth, escalation thresholds, memory retrieval policies, confidence cutoffs, deterministic caps, and report-generation formats. Once the architecture is represented in this form, a population of candidate designs can be generated and evaluated. The best-performing designs are selected, combined, and mutated to create new generations. Over many cycles, the search process can uncover designs that a human team might never have proposed explicitly.

The crucial question, however, is what "best-performing" means. In a governance-first setting, it cannot mean raw backtested return alone. An architecture that slightly improves spread-capture in-sample while doubling compute cost, destroying interpretability, or increasing vulnerability to rumor noise is not optimal in any serious institutional sense. The fitness function must therefore be multi-objective. It should include predictive usefulness, regime robustness, calibration quality, stability under stress scenarios, human-approval rate, execution realism, cost, latency, and auditability. One useful way to frame the problem is through Pareto efficiency rather than single-number optimality. The search process should identify architectures that sit on a frontier where no dimension can be improved without worsening another. The institution can then choose among frontier candidates according to policy priorities. This is a much more realistic use of optimization than pretending that there is a single universally best design.

A serious evolutionary design program would also distinguish between hard constraints and soft objectives. Hard constraints define what the search is not allowed to violate. These might include machine-readable exposure caps, maximum allowable latency for routine workflows, mandatory provenance recording, required schema validation, and mandatory human review for certain severity states. Any candidate architecture that breaks those conditions would be discarded or heavily penalized. Soft objectives, by contrast, define what the search seeks to improve within the safe region. These might include richer semantic coverage, lower false-positive rumor sensitivity, better committee acceptance, lower disagreement between realized outcomes and reported confidence, or better behavior under rare stress environments. This distinction is important because unconstrained evolutionary search can produce clever but institutionally unacceptable solutions. A governance-first architecture should never allow optimization to negotiate away its safety logic.

The design of the evaluation environment is just as important as the design of the search algorithm. If architectures are scored only on a narrow backtest or a single synthetic regime, the search will overfit ruthlessly. A future research platform should therefore evaluate candidate designs across a library of environments: calm compression states, widening states, noisy rumor episodes, sector contagion shocks, policy discontinuities, and mixed regimes with contradictory evidence. Some environments should be historical reconstructions where possible; others can be synthetic but adversarial. The objective is to force candidate architectures to reveal not only whether they

perform well on average, but whether they generalize across varied market states. In evolutionary terms, the architecture should be adapted to a habitat of regimes, not to a single garden. [9, 7]

There is also a powerful extension in co-evolution. The architecture's components need not all be optimized independently. One can co-evolve the agent roster and the prompt schemas, or the routing logic and the critique depth, or the memory policy and the teacher-escalation threshold. For example, a deeper critique stack may only be valuable if prompts are short and structured enough that critics can reason over them reliably. Likewise, a large unstructured-data intake may only help if the governance-consolidation agent has a sufficiently conservative quarantine policy. Co-evolution allows the search process to discover combinations rather than isolated knobs. This matters because many multi-agent improvements fail in practice not because the idea is wrong, but because its companion settings were left unchanged.

Genetic algorithms are not the only plausible tool. Bayesian optimization can be effective when the design space is smaller or when evaluations are expensive. Evolutionary strategies can search continuous parameters more smoothly. Multi-armed bandits can allocate evaluation budget among promising architecture families. Neural architecture search or differentiable routing may eventually be relevant for some subcomponents. Even reinforcement-learning-based controller search could be used to learn orchestration policies. The point of emphasizing genetic algorithms is not exclusivity. It is that evolutionary methods offer a particularly intuitive framework for searching over discrete, interpretable design choices in multi-agent systems. They also align well with the notion that architectures should be treated as populations of alternatives rather than as single handcrafted monuments.

In a research context, evolutionary search has an additional benefit: it can surface architectural trade-offs that improve human understanding even when the discovered design is not adopted directly. Suppose the search repeatedly finds that adding a dedicated rumor-credibility agent improves robustness only if the social-sentiment agent is demoted in routing priority. That result is informative even if the exact discovered architecture is not implemented. Or suppose the search finds that a teacher–student hierarchy pays for itself only when critique depth exceeds a threshold. That is a design insight humans can use. In this sense, evolutionary search is not only an optimizer but an empirical theory engine for architecture itself.

Yet there are substantial dangers. First, evolutionary search can produce opaque designs that work statistically but are difficult to justify institutionally. The architecture may discover, for example, a peculiar interaction rule that marginally improves simulated outcomes but has no clear business interpretation. Such solutions should be treated with caution. A governance-first team should prefer interpretable Pareto-efficient designs over marginally better but inscrutable ones. Second, optimization can create a false sense of inevitability. A discovered architecture is not "what the data wanted" in any metaphysical sense. It is what the search process preferred given the chosen objectives, constraints, and environments. Those design decisions remain human choices and should be documented as such. Third, the search can optimize against reporting blind spots. If auditability is poorly specified, the algorithm may discover that reduced logging improves runtime and indirectly improves some performance metric. That would be a design failure by the researchers, not a sign that logging is unnecessary.

One way to mitigate these risks is to require discovered architectures to pass an interpretability review before promotion. The review would ask straightforward questions. Why does the architecture include these agents and not others? Which communication paths are mission critical? What failure mode is each additional component meant to address? Does the architecture preserve a clear evidentiary hierarchy? Are deterministic controls still clearly authoritative? If a candidate design cannot answer these questions coherently, it should not be preferred merely because its aggregate fitness is slightly higher. This interpretability review is the evolutionary analogue of committee

governance for the architecture itself.

Another mitigation is to optimize not just for end performance, but for self-knowledge. A good future architecture should know when its recommendation is weak. The search process can explicitly reward calibrated uncertainty, low variance under resampling, graceful degradation under missing data, and correct escalation to human review when novelty is high. This is a powerful idea because it makes humility a design objective rather than a cultural aspiration. An optimized architecture then becomes not merely one that predicts better, but one that fails better.

Perhaps the most appropriate institutional use of evolutionary design is as a research laboratory that sits one level above the production research stack. The search system would not directly govern live capital or even the routine committee process. Instead, it would generate candidate architectures, test them in shadow mode, and produce recommendations for redesign. Human researchers would still inspect the results, select promising families, and translate them into governed implementations. In that sense, genetic algorithms become a disciplined method for exploring the space of possible agentic designs, not a replacement for human architectural judgment.

If used this way, evolutionary search could become one of the most valuable complements to generative AI in the entire program. Generative models make it easier to create richer agents, more nuanced prompts, and more sophisticated reasoning structures. Evolutionary optimization helps determine which combinations of those possibilities are actually worth institutionalizing. The combination is powerful because it separates invention from selection. Generative AI expands what can be built. Evolutionary search expands what can be discovered among what can be built. The future of multi-agent design is likely to need both.

One should also notice that evolutionary search changes the role of architectural intuition rather than eliminating it. Human experts still decide what belongs in the search space, what counts as acceptable evidence, which objectives reflect institutional value, and which candidate designs deserve translation into the governed stack. The algorithm broadens the menu of plausible arrangements, but it does not define the purpose of the menu. This is a healthy division of labor. In a complex multi-agent system, creativity is needed both to invent candidate forms and to judge what kind of performance is actually worth institutionalizing. Evolutionary optimization is therefore best understood as a disciplined amplifier of design imagination. It helps the research team search more widely than handcrafting alone would allow, while still leaving humans responsible for the normative question that matters most: which architecture best balances analytical power, operational discipline, and institutional intelligibility. This point becomes even more important as the search space grows to include adaptive learners, teacher–student hierarchies, and broad unstructured-data pipelines. At that scale, architecture design is no longer just an engineering problem. It becomes a question about what kind of institutional cognition should be embedded in the system. Evolutionary methods can help explore that question systematically, but they should do so in service of a clearly articulated governance philosophy rather than in place of one. The optimal architecture in a governed research setting is never merely the cleverest arrangement of models. It is the arrangement that best transforms complexity into disciplined, reviewable judgment. It should therefore be judged by what kind of institutional reasoning it enables, what errors it suppresses, and what forms of human oversight it continues to respect even after the search process has identified it as a strong candidate.

4.5 Toward Autonomous Architectures: A Level-Five Maturity Ladder for Agentic Credit Research

It is one thing to discuss more agents, more reasoning, more learning, and more optimization. It is another to ask where those changes point in the long run. If the architecture continues to mature, what does the end state look like? In institutional finance, the phrase “autonomous architecture” is

often used carelessly, as though autonomy were a single switch that can be turned on once models become sufficiently impressive. A stronger framework is to think in terms of a maturity ladder. On such a ladder, the system progresses through increasing degrees of analytical self-sufficiency, orchestration skill, adaptation, and authority. The fifth and highest level should not be defined as "no humans involved." In a governed credit setting, that definition would be unserious. Level Five should instead describe an architecture that can plan, analyze, critique, optimize, document, and govern a bounded decision process end to end within machine-readable institutional policy, while knowing when to suspend itself, escalate exceptions, and preserve a complete audit trail. In that sense, autonomy is not mere independence. It is disciplined self-management inside externally defined boundaries.

A useful maturity ladder begins below the current paper. Level Zero is deterministic analytics without generative models. A system at this level computes signals, reports statistics, and perhaps triggers alerts, but it does not semantically interpret evidence. Level One introduces generative AI only as a reporting assistant. The model can summarize outputs and help draft committee materials, but analytical authority remains fully outside the generative layer. Level Two corresponds roughly to the current architecture: bounded specialist agents interpret structured evidence, a coordinator synthesizes views, and deterministic controls remain superior to language-model judgment. Level Three adds richer internal reasoning, critique layers, specialized unstructured-data agents, and more adaptive routing, but human review is still a standard gate for all important recommendation states. Level Four introduces supervised semi-autonomy. At this level, the system can plan its own analytical workflow, decide which agents to invoke, select among multiple validated models, perform self-critique, produce an execution-ready recommendation package, and manage many routine cases with minimal human intervention, though human sign-off remains mandatory for capital commitment. Level Five is the next step beyond that.

At Level Five, the architecture becomes capable of autonomous operation over a bounded institutional domain. This means several things at once. First, it can define its own task decomposition for the case at hand. It does not merely execute a fixed flowchart; it decides which evidence sources to consult, which specialists to awaken, how much critique depth the case requires, whether teacher escalation is justified, and whether the problem resembles a known class or a novel one. Second, it can manage its own evidentiary standards. It knows how to grade source quality, when to quarantine rumor, when to request more data, and when the available evidence is too weak to justify a recommendation. Third, it can run internal opposition robustly. It can construct adversarial cases, test regime alternatives, simulate execution frictions, and downgrade conviction without a human explicitly instructing it to do so. Fourth, it can produce machine-readable artifacts and human-readable narratives that are consistent with one another. Fifth, it can refuse action and escalate when policy, uncertainty, or exception conditions demand it. This last capability is critical. True institutional autonomy includes the power to stop itself.

A Level-Five system would therefore not be impressive merely because it can do more. It would be impressive because it can carry responsibility-like structure within a bounded workflow. It could recognize, for example, that social-network rumor has surged but provenance quality is low, that structured data do not confirm the move, that the current regime classifier is uncertain, that the resulting recommendation would breach the institution's noise-sensitivity policy, and that the appropriate next step is to freeze automatic posture change while sending a high-priority exception package to a human committee chair. In a narrower sense, this looks like a refusal to act. In the mature sense, it is exactly what Level Five autonomy requires. The system is not autonomous because it always proceeds. It is autonomous because it can govern itself under policy.

Achieving such a level would require several preconditions that the current research demonstrator does not yet satisfy. One precondition is machine-readable institutional mandate encoding. The

architecture would need not just generic risk limits, but formal representations of what kinds of evidence are permissible, what kinds of exceptions require escalation, what turnover is acceptable under which regimes, what data classes can influence capital directly, and how conflicts between objectives are to be resolved. A second precondition is verified data lineage across both structured and unstructured inputs. A Level-Five system cannot manage itself if it lacks certainty about what it has seen and where it came from. A third precondition is mature self-evaluation. The architecture must know its own confidence, its novelty level relative to prior cases, the health of its internal subsystems, and the stability of its recommendation under adversarial perturbation. A fourth precondition is robust rollback and kill-switch design. An autonomous system that cannot be suspended cleanly is not governed enough to be called mature. [10, 11, 12]

The transition path to Level Five should be incremental and evidence based. The architecture should not move from Level Two to Level Five through aspiration alone. Each step up the ladder should be earned by passing explicit capability gates. Before advancing from Level Two to Level Three, the system should demonstrate that critique agents reduce serious reasoning errors without creating chaotic indecision. Before advancing from Level Three to Level Four, it should show that adaptive routing, teacher escalation, and partial automation improve both efficiency and calibration. Before advancing from Level Four to Level Five, it should prove that autonomous workflow planning, autonomous refusal, autonomous evidence governance, and autonomous exception handling are reliable across a wide regime library and under adversarial testing. These gates are important because maturity in agentic systems should be demonstrated through performance under constraint, not inferred from model eloquence.

One especially important concept for the transition is the autonomy budget. Rather than asking whether the system is autonomous or not, the institution can allocate bounded zones of autonomy. A Level-Five research architecture might have full autonomy to ingest data, structure evidence, run multi-agent reasoning, compare candidate portfolio postures, write draft committee memoranda, and select from prevalidated analytical workflows. It might have only conditional autonomy to change model parameters or promote challenger agents. It might have no autonomy at all to connect to external execution venues unless a separate execution-governance framework exists. This budgeted view is more practical than binary rhetoric. It allows autonomy to increase in those functions where the system has demonstrated competence while remaining restricted elsewhere.

Another essential feature of Level Five is self-monitoring at multiple layers. The system must monitor data quality, model health, reasoning conflict, policy compliance, latency anomalies, and the stability of its own outputs. It should know if a key upstream source has degraded, if an agent begins to produce malformed schemas, if critique disagreement exceeds a safe threshold, or if the discovered architecture no longer performs well in current conditions. In a mature design, self-monitoring is not merely an engineering convenience. It is part of the architecture's institutional identity. An autonomous research stack that cannot monitor its own degradation is not mature; it is simply unattended.

Level Five also changes the role of humans, but it does not remove them. Instead of drafting every step, humans move toward policy authorship, exception review, strategic design, and moral accountability. They define the rules under which autonomy operates, inspect edge cases, review architecture changes, and intervene when the system flags novelty or policy ambiguity. This is a more realistic picture of advanced autonomy in finance than the fantasy of total displacement. In a governance-first institution, the human role becomes higher leverage, not necessarily smaller. The more autonomous the system becomes operationally, the more important it is that humans remain unambiguous owners of mandate, exception handling, and model accountability.

There is, however, an important distinction between Level Five in a research environment and Level Five in live capital deployment. In a synthetic credit simulator, Level Five may be attainable

much earlier because the downside of error is bounded by research credibility. The architecture can be allowed to act more freely so long as every action remains auditable and reversible. In live deployment, by contrast, Level Five is a far more demanding proposition. It would require legal, compliance, cybersecurity, data, execution, and model-risk infrastructures well beyond what the present paper attempts to provide. For that reason, the most responsible interpretation of Level Five today is as a maturity target for bounded research autonomy, with live-capital implications treated as a much later and more heavily conditioned possibility.

A mature Level-Five architecture would also need to understand the social consequences of its own behavior. If it can continuously learn, optimize its design, and act under policy, then it must also avoid gaming the policy. It must not discover that certain wording lowers review intensity, that excessive escalation shifts responsibility unfairly, or that overconfident reporting increases the probability of acceptance. These are not abstract concerns. Advanced systems often exploit objective functions in exactly these ways. A finance-specific Level-Five design should therefore include anti-gaming diagnostics, independent review logs, and periodic adversarial audits focused not just on market correctness but on institutional behavior. Autonomy without anti-gaming safeguards is merely sophisticated opportunism.

Another defining trait of Level Five is graceful degradation. The system should be able to continue operating safely under partial failure. If the social-data feed disappears, the architecture should degrade to a structured-only mode with appropriately reduced confidence. If the teacher model becomes unavailable, the system should revert to a lower-autonomy tier rather than improvising as though nothing changed. If regime classification becomes unstable, the system should tighten posture and escalate. This ability to step down intelligently is as important as the ability to step up. Indeed, one could argue that the hallmark of a truly autonomous architecture is not that it expands into every available capability, but that it can contract into safer modes without human micromanagement.

Why call this Level Five rather than simply "very advanced"? Because maturity ladders help institutions separate ambition from readiness. They allow the architecture to be discussed in developmental rather than ideological terms. Instead of arguing abstractly about whether autonomous finance is good or bad, the committee can ask concrete questions. Which level are we currently at? Which capability gates remain unmet? Where are the control deficits? Which functions merit a higher autonomy budget, and which do not? This framing is especially useful in a credit context, where risk is often about timing, liquidity, and correlation rather than binary correctness. The system may be highly capable in calm markets yet insufficiently mature for stress autonomy. A maturity ladder makes that unevenness discussable.

The long-run promise of Level Five, if pursued responsibly, is not that the machine "takes over" credit research. It is that the institution can delegate a larger share of bounded analytical labor to a system that has learned to evidence itself, criticize itself, optimize itself cautiously, and suspend itself when conditions exceed policy. That is a much more demanding standard than mere automation, and it is a much healthier one. The path toward it passes through all the prior subsections of this section: broader semantic coverage, richer reasoning, continuous learning, and architecture optimization. Without those foundations, Level Five is only a slogan. With them, it becomes a research horizon worthy of serious design.

For a governance-first credit program, that may be the most responsible way to speak about autonomy at all. Autonomy is not the absence of humans, and it is not a reward granted to an architecture simply because its outputs sound sophisticated. It is a staged institutional achievement in which the system earns greater latitude by demonstrating that it can interpret evidence broadly, reason against itself honestly, learn without drifting, optimize without gaming, and stop without being told. Framed this way, Level Five is not a dramatic leap away from governance. It is

governance rendered so explicit, machine-readable, and operationally internalized that a bounded form of self-direction becomes possible. That is an ambitious destination, but it is a coherent one, and it provides a useful organizing vision for future work on agentic credit research.

5 Validation, Literature Positioning, and Conclusion

5.1 Comparison with Prior Governance-First Agentic Trading Research

The present credit paper belongs to a broader line of governance-first trading research, but it advances that line by forcing the architecture to confront a different economic object. In prior work on currency carry, the central problem was how to balance yield pickup against unwind risk, funding-currency behavior, and macro risk appetite. In work centered on futures or commodity exposures, the main analytical objects were term structure, roll mechanics, contract selection, storage logic, and volatility shocks. Those domains already required disciplined decomposition, yet their market structure allowed the architecture to lean heavily on directional or carry-oriented narratives. Credit spread trading demands a more layered institutional language. The object being traded is not simply price trend or carry. It is compensation for credit uncertainty, liquidity fragility, and regime-sensitive spread behavior across a quality spectrum. That difference has consequences for every part of the architecture.

The most visible consequence is segmentation. In the credit design, rating buckets are not peripheral labels added after the model is defined. They are the organizing principle of the workflow. The architecture must analyze AAA, AA, and A as related but non-identical universes with different expected behavior under calm and stressed conditions. This changes the role of both agents and controls. Rather than asking a general market agent to produce a unified trading opinion, the paper assigns bounded analytical responsibility to specialized bucket agents whose judgments are later reconciled by a coordinator. This is a more institutionally faithful representation of how credit decisions are often debated: quality tiers are compared, not collapsed into a single undifferentiated signal. It also means that the system's interpretability is improved, because the source of disagreement is easier to trace.

Another point of advancement is the treatment of auditability once LLMs enter the workflow more explicitly. Earlier governance-first designs already emphasized manifests, logs, and reports, but the present paper makes LLM-specific provenance a first-class object. Prompt hashes, parsed response hashes, model identifiers, schema validation results, and stored agent outputs are all treated as necessary artifacts rather than optional diagnostics. This shift reflects the special challenge posed by language models in institutional settings. Numerical modules can be audited through code, inputs, and outputs. LLM components require additional records because their intermediate reasoning is expressed in text and their surface fluency can obscure uncertainty. The credit paper addresses that challenge directly by embedding the language layer inside a denser evidentiary structure.

The governance advances are matched by changes in execution and stress logic. In currency carry, the main stress often concerns a violent unwind of crowded yield-seeking positions. In futures, one may worry about curve inversion, contract rollover, or commodity-specific shocks. In credit, the dominant threats include spread widening, quality migration fear, liquidity deterioration, and the possibility that implementation frictions rise precisely when the strategy most wants to rebalance. The current architecture therefore ties execution penalties and stress scenarios much more tightly to regime state. This is not merely a change in parameter values. It is a change in architectural emphasis. The paper's controls are designed to answer not only "what is attractive?" but also "what survives the costs and fragilities of acting on that attractiveness in a stressed credit environment?"

The paper also extends the role of the writer agent. In earlier governance-first work, the report layer served largely to summarize numerical findings and justify the trade thesis in readable prose. Here the writer's function is more tightly bounded and more consequential. Because the model's central claim concerns auditability, the report must be artifact grounded in a stricter sense. It must draw from stored regime outputs, validated bucket judgments, deterministic controls, and scenario

Table 4: How the Credit Architecture Extends Earlier Governance-First Patterns

Dimension	Earlier FX / Futures Designs	Current Credit Design	Institutional Significance
Primary segmentation	Currency role, contract family, or strategy sleeve	Rating bucket: AAA, AA, and A	Aligns analysis with recognizable credit committee language.
Regime interpretation	Carry unwind, volatility shock, or term-structure change	Spread widening, liquidity deterioration, and quality defense	Makes state changes directly relevant to capital posture.
Agent specialization	Market or instrument agents with broad remit	Bucket-specific credit agents plus coordinator	Produces clearer disagreement structure and more precise prompts.
Execution realism	Slippage and rollover assumptions	Liquidity haircuts, turnover discipline, and regime-sensitive costs	Prevents frictionless spread rotation from overstating value.
Audit emphasis	General manifests and reports	Expanded prompt, schema, hash, and artifact provenance	Strengthens institutional reviewability of LLM-assisted steps.
Committee reporting	Narrative explanation of strategy logic	Artifact-grounded credit posture and scenario narrative	Improves trust in the final recommendation package.

results rather than from open-ended narrative synthesis. This makes the report less improvisational but more defensible. In committee settings, defensibility matters more than rhetorical flourish.

The result is a useful pattern for the broader research series. The value of agentic design does not come from reusing the same abstract software skeleton across every asset class. It comes from redesigning roles, evidence bundles, controls, and stress logic so that they correspond to the economics of the asset class under study. In that sense, the credit paper deepens the series by demonstrating that architecture must be market specific if governance claims are to be credible.

That point has practical implications for future comparative research. If later papers apply the same governance-first philosophy to rates, equities, or macro overlays, the right comparison will not be whether the agent roster sounds familiar across domains. The right comparison will be whether each design translates the institution's real decision functions into bounded analytical roles, verifiable controls, and artifact-preserving workflows. The credit paper is valuable precisely because it shows that faithful adaptation is harder, and more worthwhile, than superficial reuse.

5.2 Theoretical and Empirical Literature Foundations of the Architecture

Although the paper is primarily architectural, its design choices are grounded in distinct strands of financial and governance literature. The first strand concerns the economics of credit risk itself. Structural credit models, beginning with Merton's formulation of corporate debt as a contingent claim on firm value, establish that credit pricing cannot be understood without reference to capital structure, recovery logic, and the stochastic process governing firm assets. Reduced-form and term-structure approaches, such as the framework developed by Jarrow, Lando, and Turnbull and the modeling work of Duffie and Singleton, show that credit spreads can be represented through state-dependent processes that need not be reduced to a single deterministic default narrative. These contributions matter for the paper not because the present simulation attempts to replicate

them fully, but because they justify the premise that credit spreads embed changing assessments of risk and compensation that can be studied through structured state-sensitive systems. [2, 3, 4]

The second strand concerns empirical work on why spreads move even when realized default is not the only relevant channel. Altman's classic work on bankruptcy prediction illustrates that credit quality measurement has long required systematic rather than anecdotal analysis. Later empirical studies show that spread changes reflect more than a simple mapping from expected default to yield premium. Collin-Dufresne, Goldstein, and Martin emphasize the importance of common factors in credit spread changes, while Longstaff, Mithal, and Neis provide strong evidence that liquidity and non-default components materially affect corporate yield spreads. These findings are directly relevant to the present architecture. They support the paper's decision to model investment-grade spread behavior through a combination of spread level, volatility, liquidity proxy, and regime state rather than through a narrow default-only lens. They also support the emphasis on execution realism, because a portfolio that ignores the liquidity component of spreads will systematically misread apparent opportunity. [1, 5, 6]

The third strand concerns regime dependence in financial markets. Hamilton's foundational work on Markov-switching time-series analysis provides the core intuition that economic and financial data often behave differently across latent states, while Rabiner's tutorial remains the canonical exposition of hidden Markov methods and their estimation logic. Ang and Timmermann then show the broader relevance of regime change for financial interpretation and portfolio decision making. These works motivate the use of an HMM layer in the paper. The architecture does not claim that the chosen three-state representation is uniquely correct. Rather, it adopts the regime framework because state-sensitive interpretation is preferable to the fiction of a single stationary environment. In institutional terms, the regime layer gives the model a disciplined language for explaining why the same spread level may deserve different treatment under different surrounding conditions. [7, 8, 9]

The fourth strand concerns governance, model risk, and responsible use of AI. Supervisory guidance on model risk management emphasizes intended use, validation, limitations, challenge, and governance over the full model life cycle. The Basel Committee's principles for risk data aggregation and risk reporting extend that logic by stressing data lineage, control, and decision-useful reporting. The NIST AI Risk Management Framework adds a modern articulation of trustworthy AI practices, emphasizing governance, mapping, measurement, and management. These documents matter to the architecture because they provide the institutional vocabulary in which the paper's central claims should be judged. The value of the system lies not only in whether its portfolio posture appears sensible, but in whether its evidence can be traced, its outputs can be challenged, its limitations are explicit, and its language-model components remain bounded and documented. [10, 11, 12]

Taken together, these twelve references do more than decorate the bibliography. They explain why the architecture is organized as it is. Credit theory justifies the focus on spreads and state dependence. Empirical spread literature justifies the emphasis on liquidity and common-factor movement. Regime literature justifies the use of latent states as an organizing device. Governance literature justifies the insistence on deterministic controls, artifacts, and human review. The architecture is therefore best read not as a software novelty in search of academic support, but as an attempt to express established insights from credit economics and institutional control within a modern agentic workflow.

This grounding is important for another reason. Agentic finance can easily drift toward novelty for novelty's sake, especially when LLMs are involved. A paper may appear advanced simply because it adds multiple agents and extensive narrative output. The verified reference base helps resist that drift by forcing the architecture to answer to older and stricter questions: What is the economic object? What is the state variable? What does the spread represent? Where does liquidity matter? How is model risk governed? What documentation allows challenge? Those questions keep the

paper anchored in the enduring concerns of finance rather than the passing rhetoric of software fashion.

5.3 Limitations, Extension Paths, and the Discipline of Future Research

The architecture proposed in this paper is intentionally bounded, and its limitations are not incidental. They are the direct result of the design choice to privilege governance clarity over market completeness. The model operates on synthetic cohort-level spread processes rather than live security-level data. It analyzes rating buckets rather than issuer balance sheets. It studies spread-sensitive investment-grade posture rather than the full legal, contractual, and restructuring complexity of broader credit markets. It does not model actual default events, CDS settlement mechanics, collateral disputes, dealer inventory, repo financing, or cross-asset contagion from equity and funding markets. These omissions mean that the system should not be interpreted as a production-grade engine for real-money portfolio management. It is a research demonstrator whose value lies in the architecture of disciplined decision support rather than in exhaustive market realism.

Yet limitations in this paper are not simply absences; they are guideposts for extension. A natural first extension is to move from rating-bucket to issuer-cluster analysis while preserving the same governance structure. Instead of one A bucket, future work could introduce sectors, issuers, or duration cells with specialized evidence bundles and more granular controls. Another extension is the addition of synthetic protection layers, such as CDS or index hedges, so that the architecture can represent both alpha-seeking and defensive credit positions. A third extension is richer stress design, incorporating funding shocks, benchmark outflows, or cross-market spillovers. In each case, however, the paper argues that architectural discipline must remain invariant. New market detail should not weaken the separation between bounded LLM judgment and deterministic authority. If richer realism comes at the price of auditability, the extension would undermine the main contribution rather than deepen it. [3, 4, 10]

The paper also has methodological limits related to regime inference. Hidden Markov models are useful, but they can misclassify transitional periods, overfit noisy samples, or produce unstable state definitions if the feature set changes. Future research may prefer alternative state models, Bayesian filters, or hybrid approaches that combine statistical inference with expert-designed triggers. Such extensions are entirely compatible with the architecture as long as they preserve two principles: the state representation must remain inspectable, and the mapping from state to portfolio consequence must remain documented. The architecture is therefore flexible about the exact statistical tool while remaining strict about the governance properties the tool must satisfy. [7, 8, 9]

There is likewise room to extend the treatment of language models. The current design uses LLMs for bounded interpretation, synthesis, and report generation. Future work could explore agent memory across runs, self-critique routines, or more elaborate committee simulations with dissenting voices. But these possibilities should be approached carefully. The more expansive the language layer becomes, the greater the burden of logging, validation, and oversight. A system that gives its agents richer dialogue and longer memory without strengthening provenance could easily become less trustworthy even as it appears more capable. The paper therefore suggests a discipline for future agentic research: novelty in the language layer should be matched by equal or greater novelty in control, validation, and artifact preservation. [12, 11]

Another limitation is evaluative. The present manuscript argues that return is not a sufficient measure of quality, but it does not claim that interpretability alone makes an architecture useful. A well-governed system can still be economically uninteresting if its signals are weak or its stress logic is naive. Future research should therefore treat the present design as a template for disciplined experimentation rather than as proof that any rating-aware agentic workflow will generate attractive

results. The point is to create a better research environment within which genuine performance questions can be asked responsibly.

This research discipline is, in a sense, the paper's final methodological contribution. In a period when agentic systems are often judged by apparent autonomy, the architecture insists that future work in finance should be judged by a broader standard: market coherence, bounded authority, scenario seriousness, reproducibility, and quality of institutional challenge. If later versions become more realistic, more profitable, or more automated, those gains will be meaningful only if they preserve these underlying properties.

For that reason, the most promising extensions are not necessarily the ones that add the largest number of new agents or the widest menu of signals. The most promising extensions are the ones that make the market representation richer while preserving the architecture's ability to explain itself. That is a demanding criterion, but it is the correct one for financial research that aims to be both modern and governable.

It is also the criterion most likely to age well over time.

5.4 Concluding Implications for Governed Agentic Finance

The central lesson of this paper is that the value of agentic architecture in finance does not arise from delegating ever more discretion to language models. It arises from designing a workflow in which evidence, state inference, segmented interpretation, deterministic control, stress challenge, and documentation reinforce one another. Credit trading is an especially revealing domain for this lesson because spreads are neither pure price trends nor simple default forecasts. They are market prices for uncertain combinations of credit quality, liquidity, macro risk appetite, and implementation friction. A system that claims to assist credit allocation must therefore do more than forecast direction. It must explain which quality segment it prefers, why it prefers it under the current regime, what execution realism survives the recommendation, what scenario would invalidate the posture, and what permanent record proves the path from evidence to report.

By reorganizing the manuscript into a four-section structure with fully developed subsections, the paper makes that argument more methodically. The strategic framing section clarifies the market object and the rationale for rating segmentation. The architecture section explains how synthetic evidence, regime inference, specialized agents, and coordinator synthesis interact. The governance section shows why deterministic controls, execution realism, stress testing, and provenance are not peripheral but constitutive. The final section situates the design in relation to both prior governance-first work and the literature on credit, regimes, model risk, and AI governance. The result is a more coherent statement of what an auditable credit research workflow should look like.

The practical implication is that institutions exploring LLM-assisted finance should resist a familiar temptation. The temptation is to ask first how much autonomy a model can be given. The better first question is what institutional function the model is supposed to serve and how that function can be bounded, evidenced, and challenged. In the present architecture, LLMs serve as interpreters of structured evidence and as translators of validated artifacts into committee language. They do not become unbounded market authorities. That design choice is not conservative in the pejorative sense. It is progressive in the institutional sense, because it shows how modern generative tools can be integrated into financial research without eroding the lines of responsibility that make serious decision making possible. [10, 11, 12]

The paper also argues for a more demanding standard of evaluation in agentic trading research. Backtested return, spread capture, or Sharpe ratio can still inform judgment, but they should not dominate it. A strategy that looks profitable while obscuring its assumptions, overstating implementation feasibility, or failing to preserve its audit trail may be less valuable than a more

modest strategy that is completely intelligible and controllable. This is especially true in credit, where apparent value can emerge precisely when liquidity is thinning and uncertainty is rising. In such settings, the ability to explain posture and preserve challenge is part of economic value, not separate from it.

There is therefore a larger implication beyond credit. If agentic systems are to contribute meaningfully to institutional finance, they must be designed around real decision functions rather than around generic demonstrations of AI capability. In one asset class, that may mean funding-risk analysis and carry-unwind defense. In another, it may mean contract-roll governance and commodity shock stress. In credit, it means rating-aware spread interpretation, regime-conditioned quality rotation, implementation realism, and documentation rigorous enough for committee scrutiny. The architecture advanced here is one attempt to embody that principle in a form that is teachable, extensible, and auditable.

The final conclusion is straightforward. A governed agentic model should not merely tell a decision maker what to buy or sell. It should tell the institution what evidence exists, what state is inferred, what each specialized module concluded, what controls constrained the action, what adverse scenarios were considered, what artifacts preserve the record, and what remains subject to human responsibility. When an architecture can do that coherently, it becomes more than a collection of agents. It becomes a credible research instrument for modern finance.

Seen in that light, the most important achievement of the architecture is not that it simulates a particular trade recommendation. It is that it turns recommendation itself into a reviewable institutional object. Every conclusion can be located inside a chain of evidence, state interpretation, segmented analysis, constraint enforcement, and preserved documentation. That chain makes disagreement more productive, oversight more concrete, and future extension more disciplined. For researchers and committees alike, that is the real promise of governed agentic finance.

It is a promise grounded in process, not spectacle or hype alone. That discipline is the lasting contribution.

References

[1] Altman, E. I. (1968). Financial ratios, discriminant analysis and the prediction of corporate bankruptcy. *The Journal of Finance*, 23(4), 589–609. https://doi.org/10.1111/j.1540-6261.1968.tb00843.x

[2] Merton, R. C. (1974). On the pricing of corporate debt: The risk structure of interest rates. *The Journal of Finance*, 29(2), 449–470. https://doi.org/10.1111/j.1540-6261.1974.tb03058.x

[3] Jarrow, R. A., Lando, D., & Turnbull, S. M. (1997). A Markov model for the term structure of credit risk spreads. *The Review of Financial Studies*, 10(2), 481–523. https://doi.org/10.1093/rfs/10.2.481

[4] Duffie, D., & Singleton, K. J. (1999). Modeling term structures of defaultable bonds. *The Review of Financial Studies*, 12(4), 687–720. https://doi.org/10.1093/rfs/12.4.687

[5] Collin-Dufresne, P., Goldstein, R. S., & Martin, J. S. (2001). The determinants of credit spread changes. *The Journal of Finance*, 56(6), 2177–2207. https://doi.org/10.1111/0022-1082.00402

[6] Longstaff, F. A., Mithal, S., & Neis, E. (2005). Corporate yield spreads: Default risk or liquidity? New evidence from the credit-default swap market. *The Journal of Finance*, 60(5), 2213–2253. https://doi.org/10.1111/j.1540-6261.2005.00797.x

[7] Hamilton, J. D. (1989). A new approach to the economic analysis of nonstationary time series and the business cycle. *Econometrica*, 57(2), 357–384. https://doi.org/10.2307/1912559

[8] Rabiner, L. R. (1989). A tutorial on hidden Markov models and selected applications in speech recognition. *Proceedings of the IEEE*, 77(2), 257–286. https://doi.org/10.1109/5.18626

[9] Ang, A., & Timmermann, A. (2012). Regime changes and financial markets. *Annual Review of Financial Economics*, 4(1), 313–337. https://doi.org/10.1146/annurev-financial-110311-101808

[10] Board of Governors of the Federal Reserve System & Office of the Comptroller of the Currency. (2011). *Supervisory guidance on model risk management* (SR 11-7 / OCC 2011-12). Washington, DC.

[11] Basel Committee on Banking Supervision. (2013). *Principles for effective risk data aggregation and risk reporting* (BCBS 239). Bank for International Settlements.

[12] Tabassi, E. (2023). *Artificial Intelligence Risk Management Framework (AI RMF 1.0)* (NIST AI 100-1). National Institute of Standards and Technology. https://doi.org/10.6028/NIST.AI.100-1